



Data-Intensive Distributed Computing

CS 451/651 (Fall 2018)

Part I: MapReduce Algorithm Design (3/4)
September 13, 2018

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

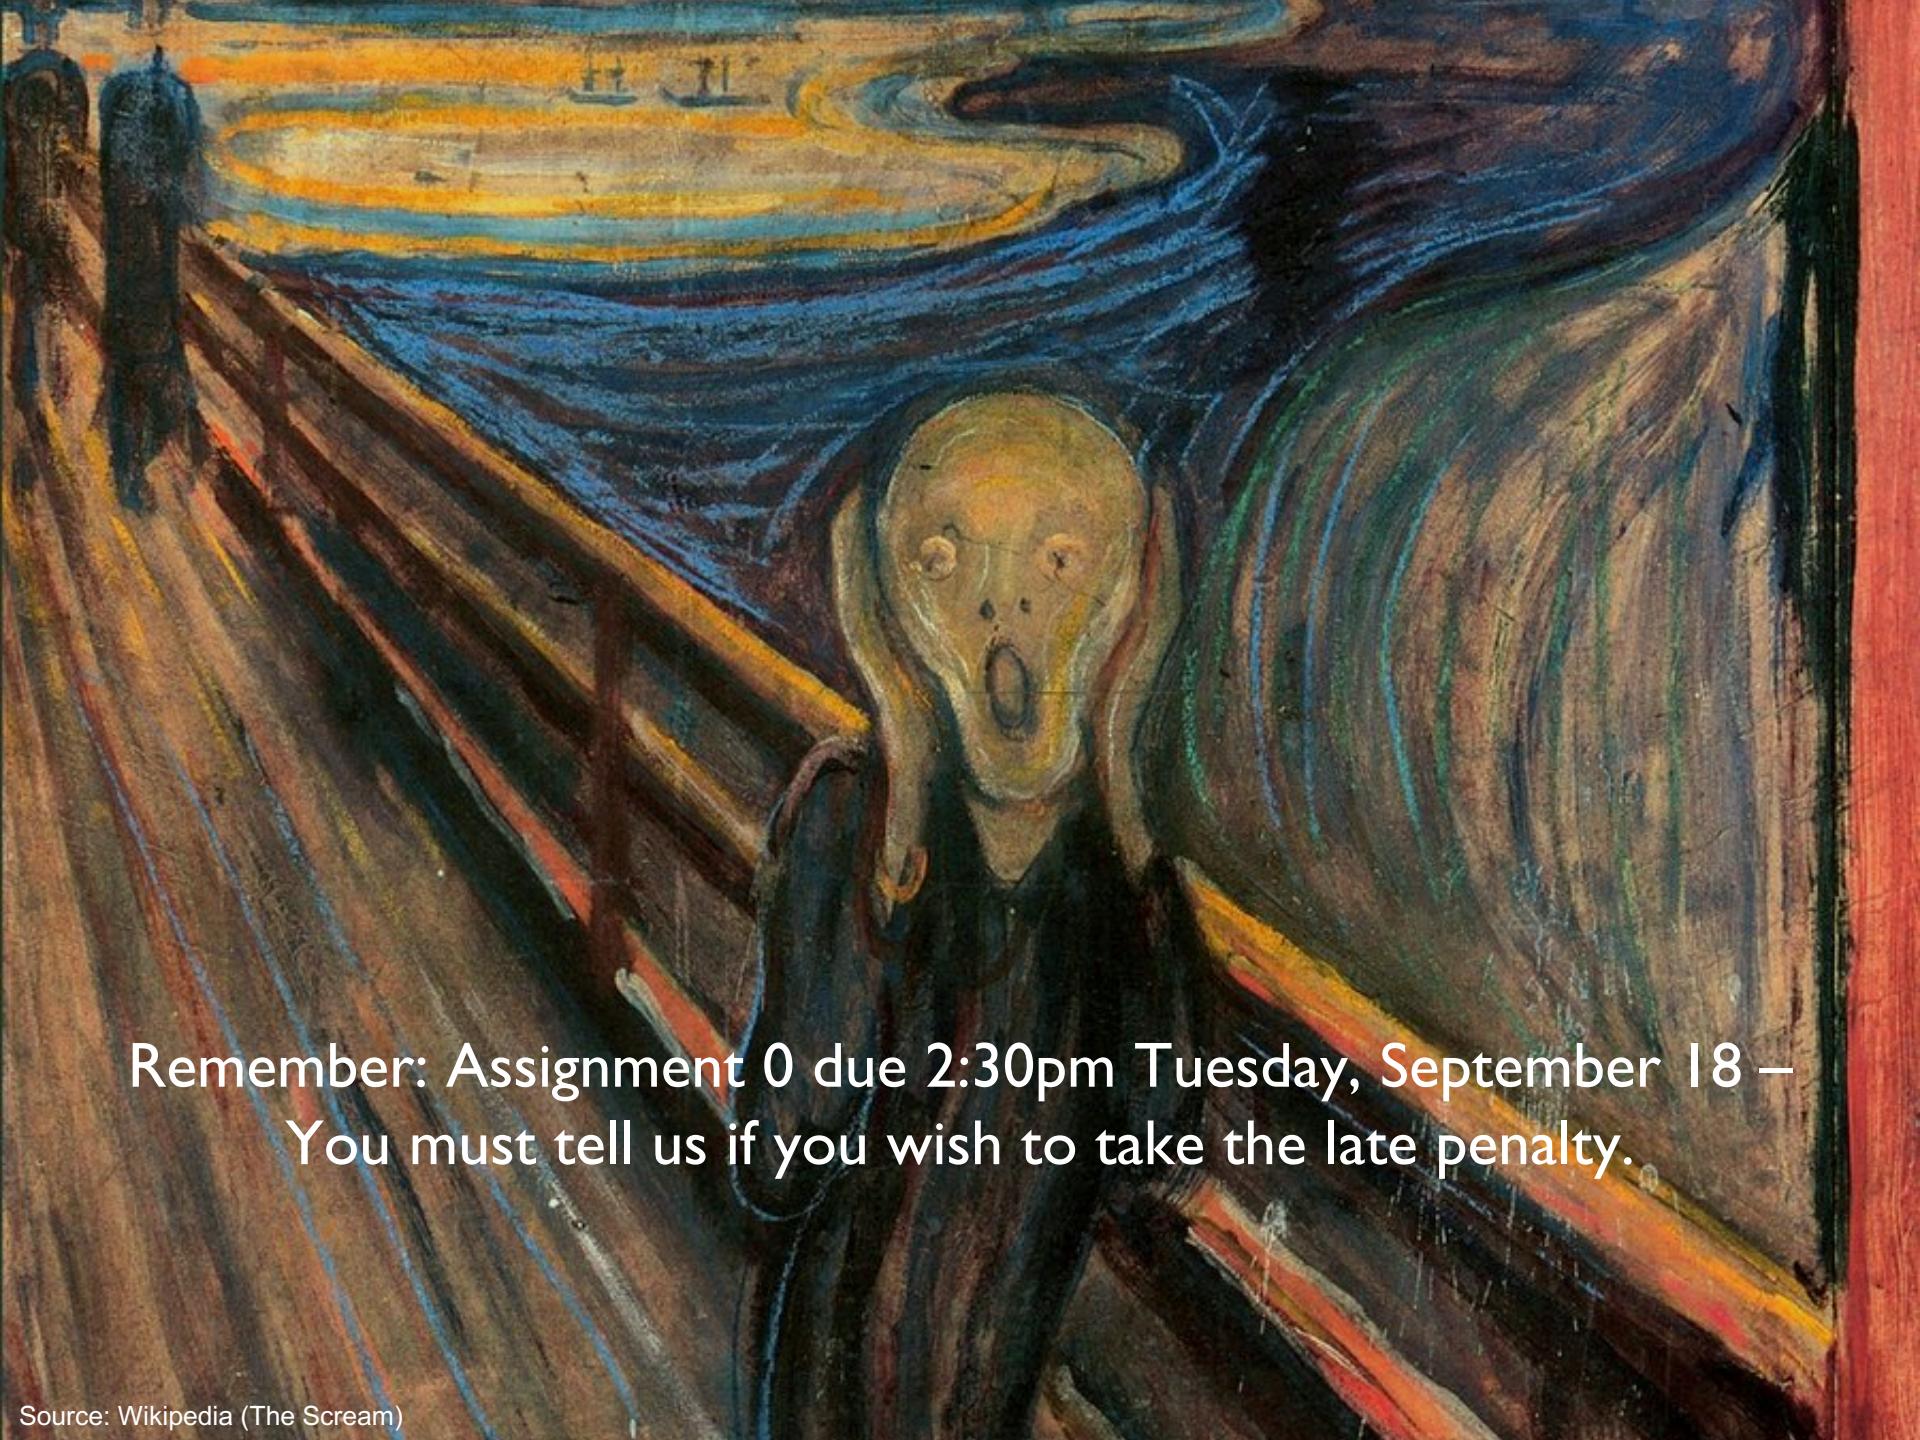
These slides are available at <http://lintool.github.io/bigdata-2018f/>



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

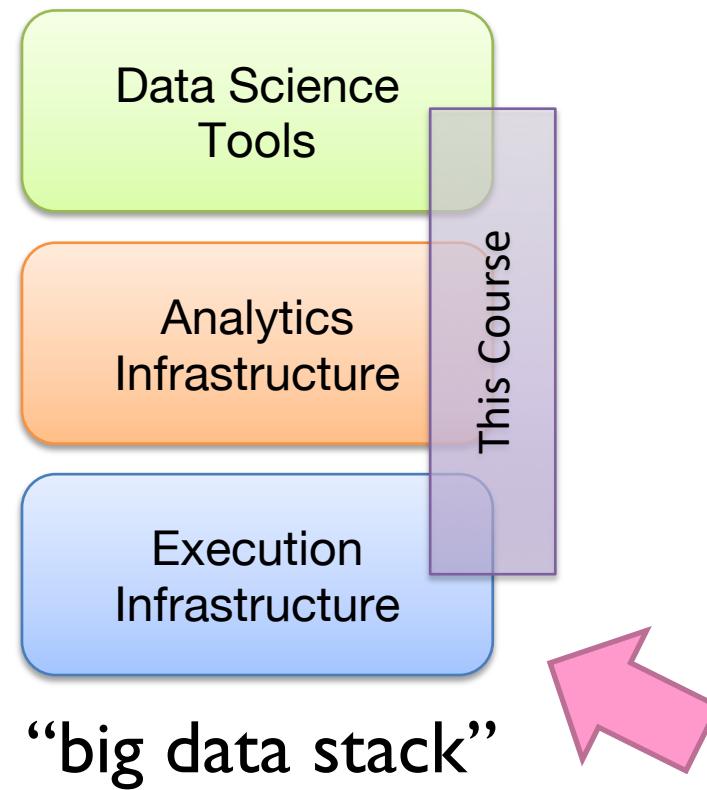
Agenda for Today

Cloud computing
Datacenter architectures
Hadoop cluster architecture
MapReduce physical execution



Remember: Assignment 0 due 2:30pm Tuesday, September 18 –
You must tell us if you wish to take the late penalty.

Today



An aerial photograph showing a vast expanse of white and grey clouds against a clear blue sky. The clouds are dense and layered, creating a textured pattern across the frame. In the lower right quadrant, a dark, mountainous landmass is visible, partially obscured by clouds.

Aside: Cloud Computing

The best thing since sliced bread?

Before clouds...

Grids

Connection machine

Vector supercomputers

...

Cloud computing means many different things:

Big data

Rebranding of web 2.0

Utility computing

Everything as a service

Rebranding of web 2.0

Rich, interactive web applications

Clouds refer to the servers that run them

Javascript! (ugh)

Examples: Facebook, YouTube, Gmail, ...

“The network is the computer”: take two

User data is stored “in the clouds”

Rise of the tablets, smartphones, etc. (“thin clients”)

Browser is the OS

GENERAL  ELECTRIC

R 13%

8 9 0 1 2 2 1 0 9 8 8 9 0 2 1 0 9 8 8 9 0 1 2
7 6 5 4 3 3 4 5 7 7 6 5 4 3 3 4 5 6 7 7 6 5 3
K I L O W A T T H O U R S

CL 200

TYPE I-60-S
SINGLE STATOR  FM 2S
WATTHOUR METER

TA 30

240V

3W

CAT. NO.

720X1G1

Kh 7.2
60~

397128

• 44 617 187 •

MADE IN U.S.A.

P
G
E
and

Utility Computing

What?

Computing resources as a metered service (“pay as you go”)

Why?

Cost: capital vs. operating expenses

Scalability: “infinite” capacity

Elasticity: scale up or down on demand

Does it make sense?

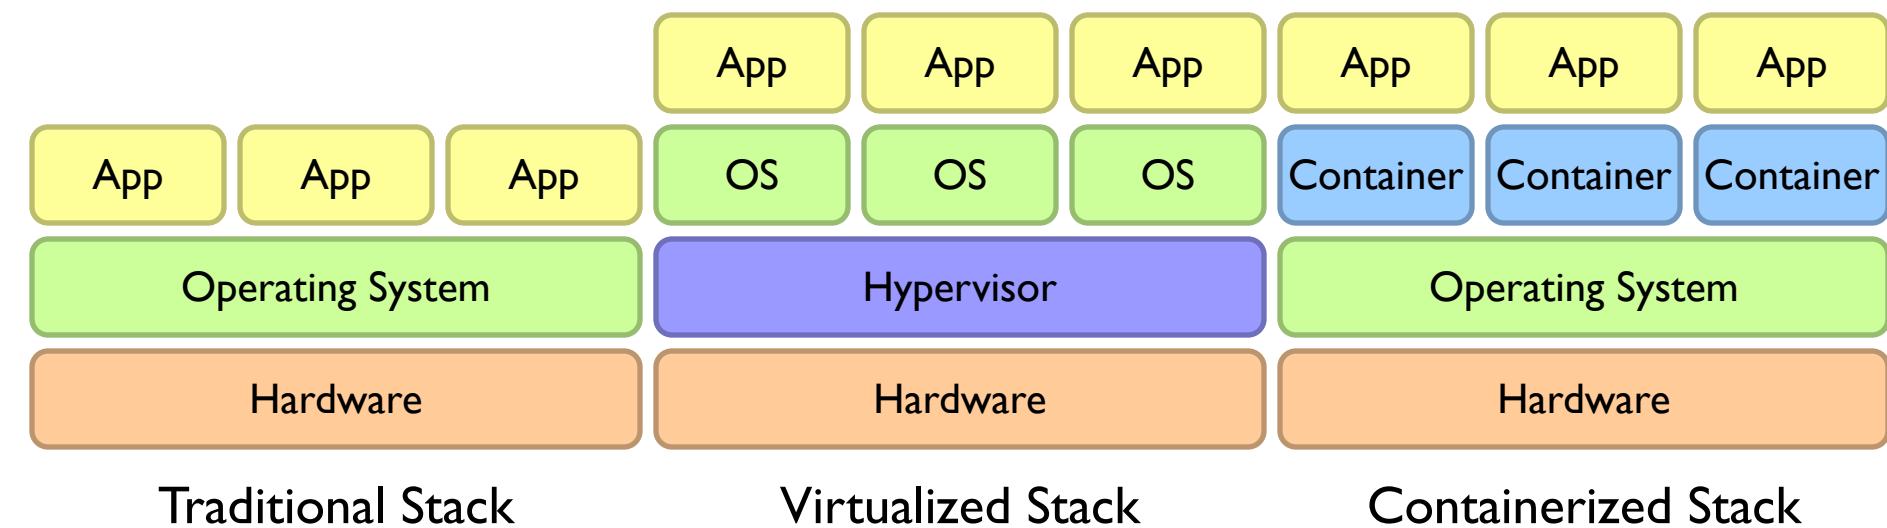
Benefits to cloud users

Business case for cloud providers

I think there is a world market for about five computers.



Evolution of the Stack



Everything as a Service

Infrastructure as a Service (IaaS)

Why buy machines when you can rent them instead?

Examples: Amazon EC2, Microsoft Azure, Google Compute

Platform as a Service (PaaS)

Give me a nice platform and take care of maintenance, upgrades, ...

Example: Google App Engine

Software as a Service (SaaS)

Just run the application for me!

Example: Gmail, Salesforce

Everything as a Service

Database as a Service

Run a database for me

Examples: Amazon RDS, Microsoft Azure SQL

Search as a Service

Run a search engine for me

Example: Amazon Elasticsearch Service

Function as a Service

Run this function for me

Example: Amazon Lambda, Google Cloud Functions

Who cares?

A source of problems...

Cloud-based services generate big data

Clouds make it easier to start companies that generate big data

As well as a solution...

Ability to provision clusters on-demand in the cloud

Commoditization and democratization of big data capabilities

An aerial photograph showing a vast expanse of white and grey clouds against a clear blue sky. The clouds are layered and textured, creating a sense of depth. In the lower right corner, a dark, mountainous landmass is visible, partially obscured by clouds.

So, what *is* the cloud?

What is the Matrix?

The Matrix is a science fiction film series directed by the Wachowskis. It consists of three main films: "The Matrix" (1999), "The Matrix Reloaded" (2003), and "The Matrix Revolutions" (2003). The story follows Neo, a computer hacker who discovers he is actually a slave in a simulated reality called the Matrix. He teams up with Trinity and Neo to fight against the machines that created the Matrix.

The Matrix is set in a futuristic world where humans are enslaved in a simulated reality. The machines have taken over the world and are using the energy produced by the human population to power their own civilization. The humans are kept in a dream-like state, unaware of their true reality. Neo, a computer hacker, begins to suspect that something is not right. He meets Trinity, another member of the resistance, and together they begin to uncover the truth about the Matrix.

The Matrix features a cast of talented actors including Keanu Reeves as Neo, Laurence Fishburne as Agent Smith, and Carrie-Anne Moss as Trinity. The film's special effects are groundbreaking, creating a convincing simulated reality that blurs the line between fact and fiction. The music score by Hans Zimmer is also highly regarded, providing a fitting soundtrack for the film's intense action scenes.

The Matrix has become a cultural icon, inspiring numerous books, articles, and even a spin-off television series. Its themes of reality vs. simulation, free will vs. determinism, and the nature of existence continue to captivate audiences around the world.





Source: Wikipedia (The Dalles, Oregon)

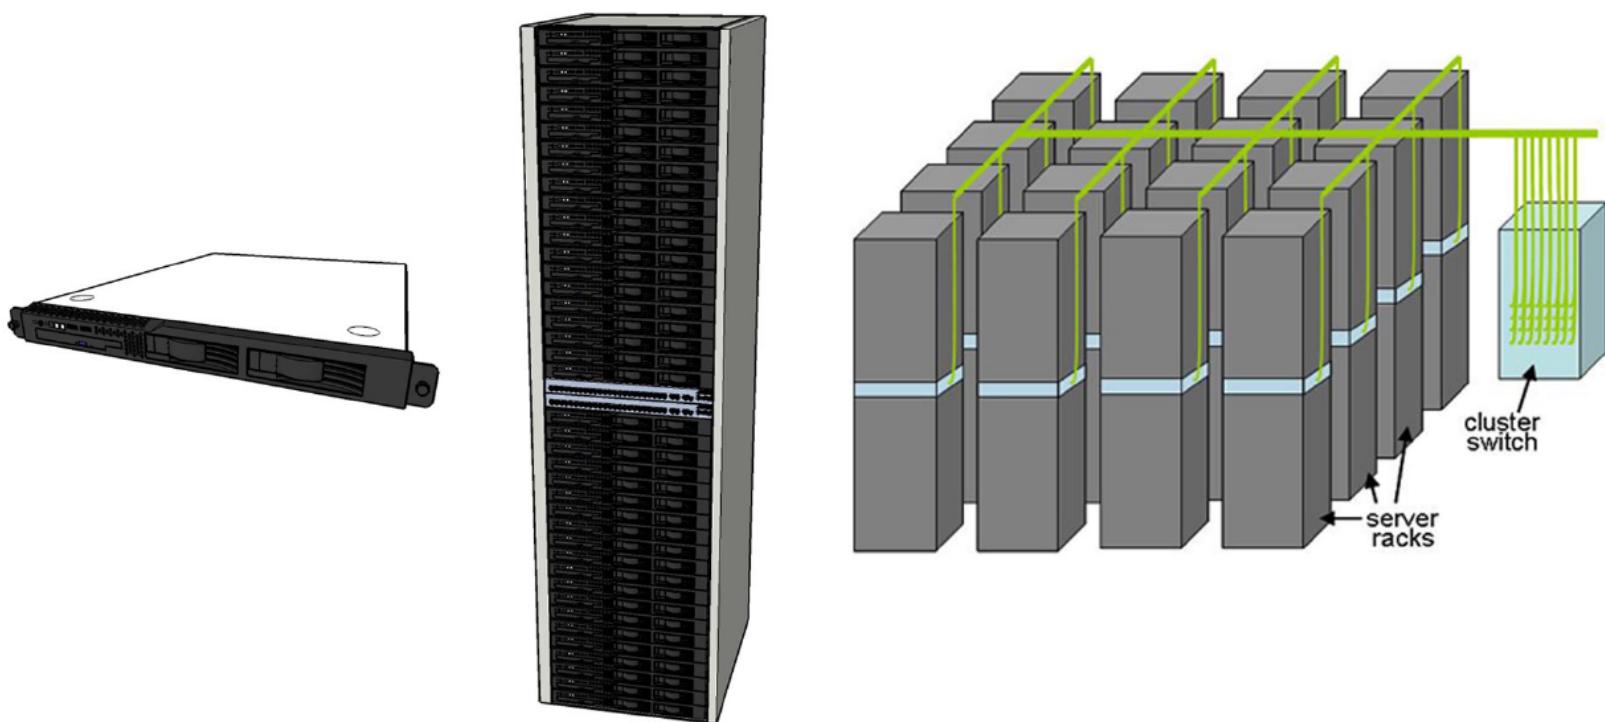


Source: Bonneville Power Administration





Building Blocks

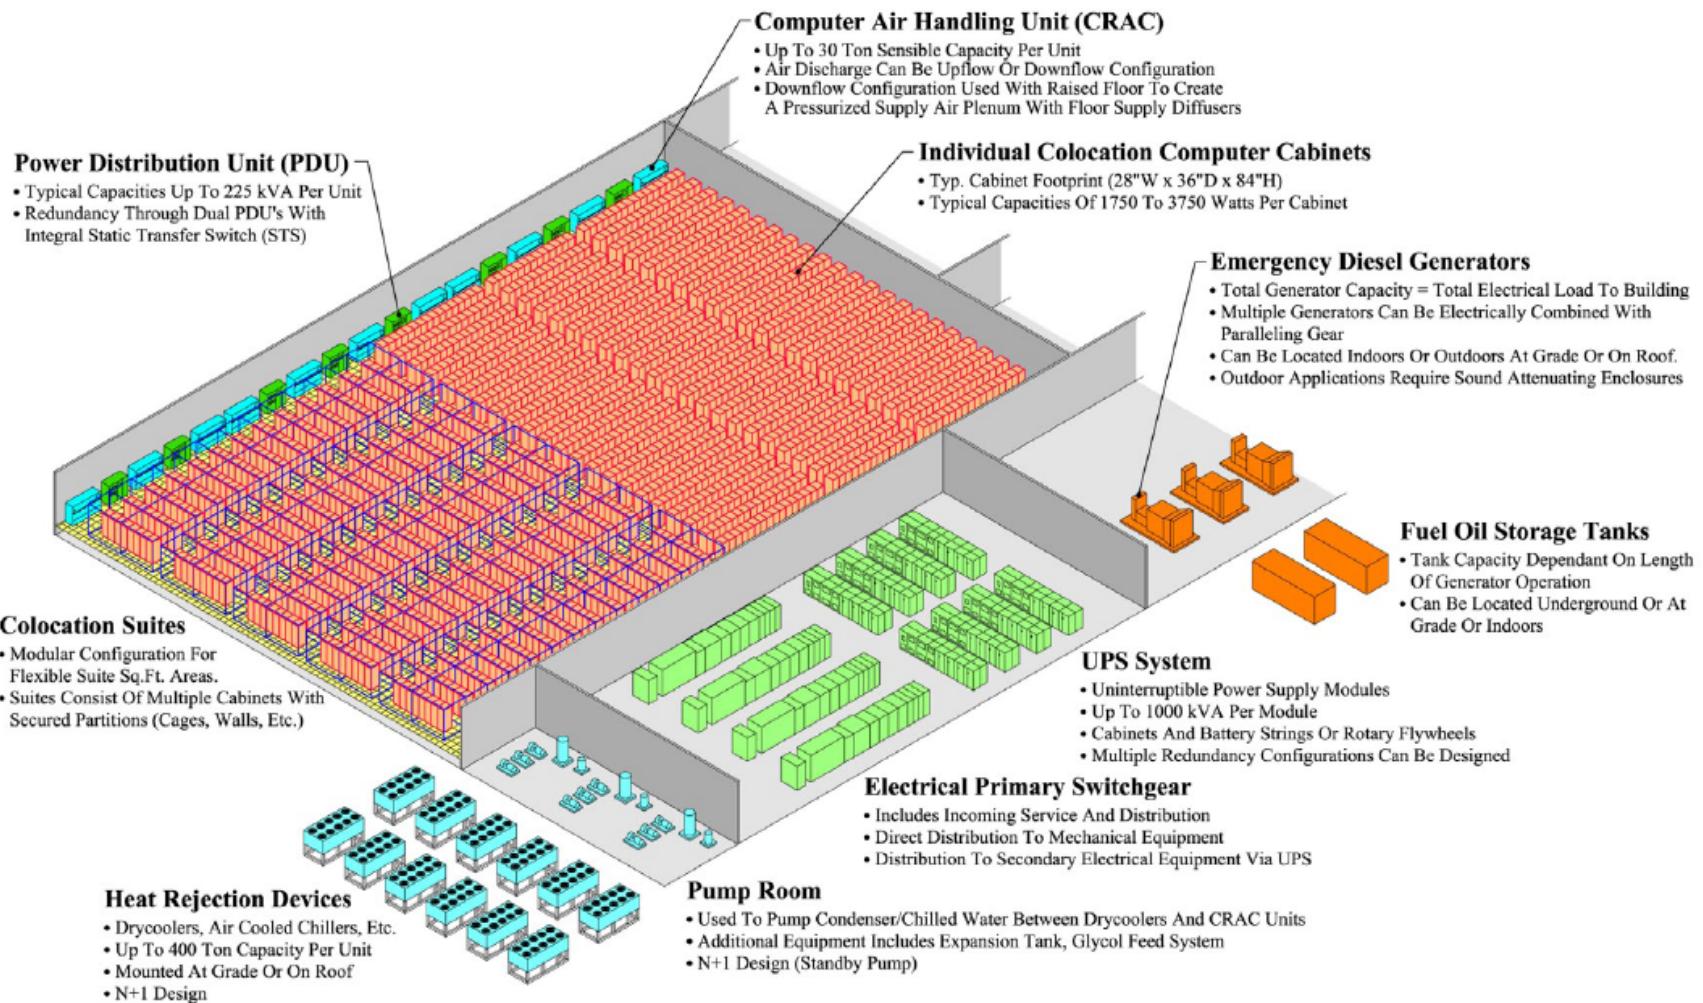






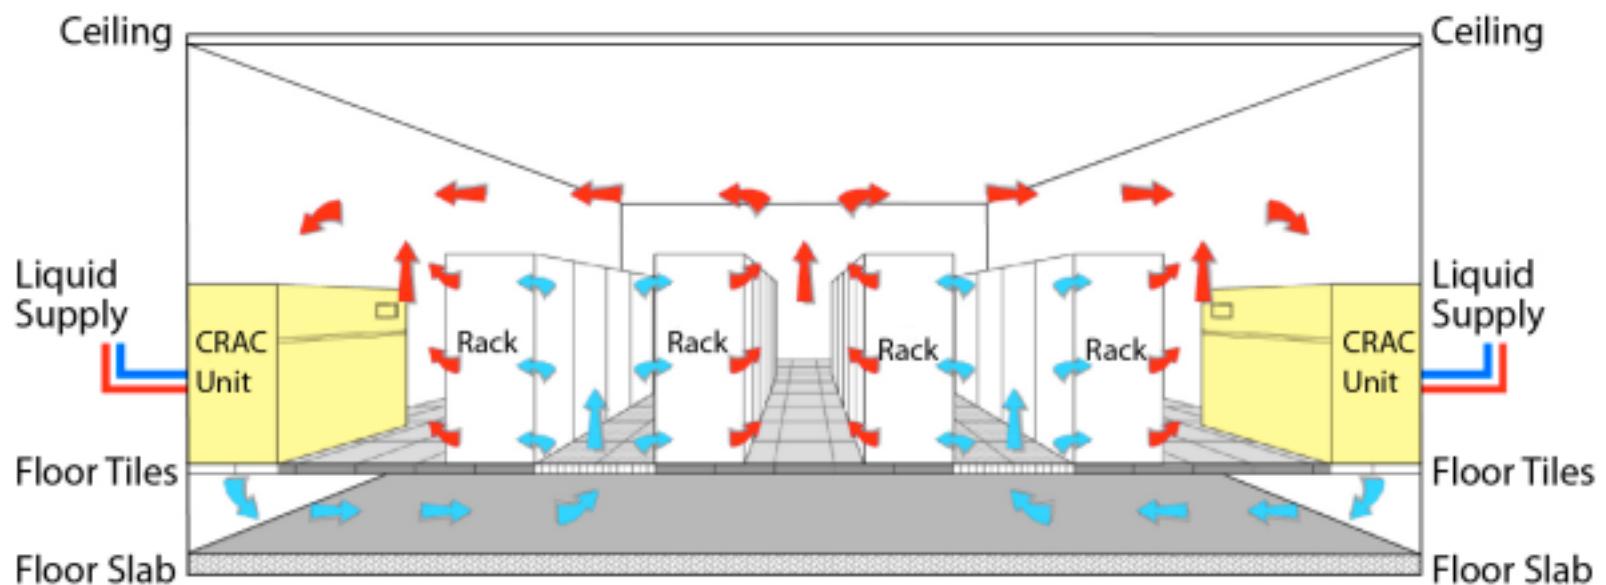


Anatomy of a Datacenter

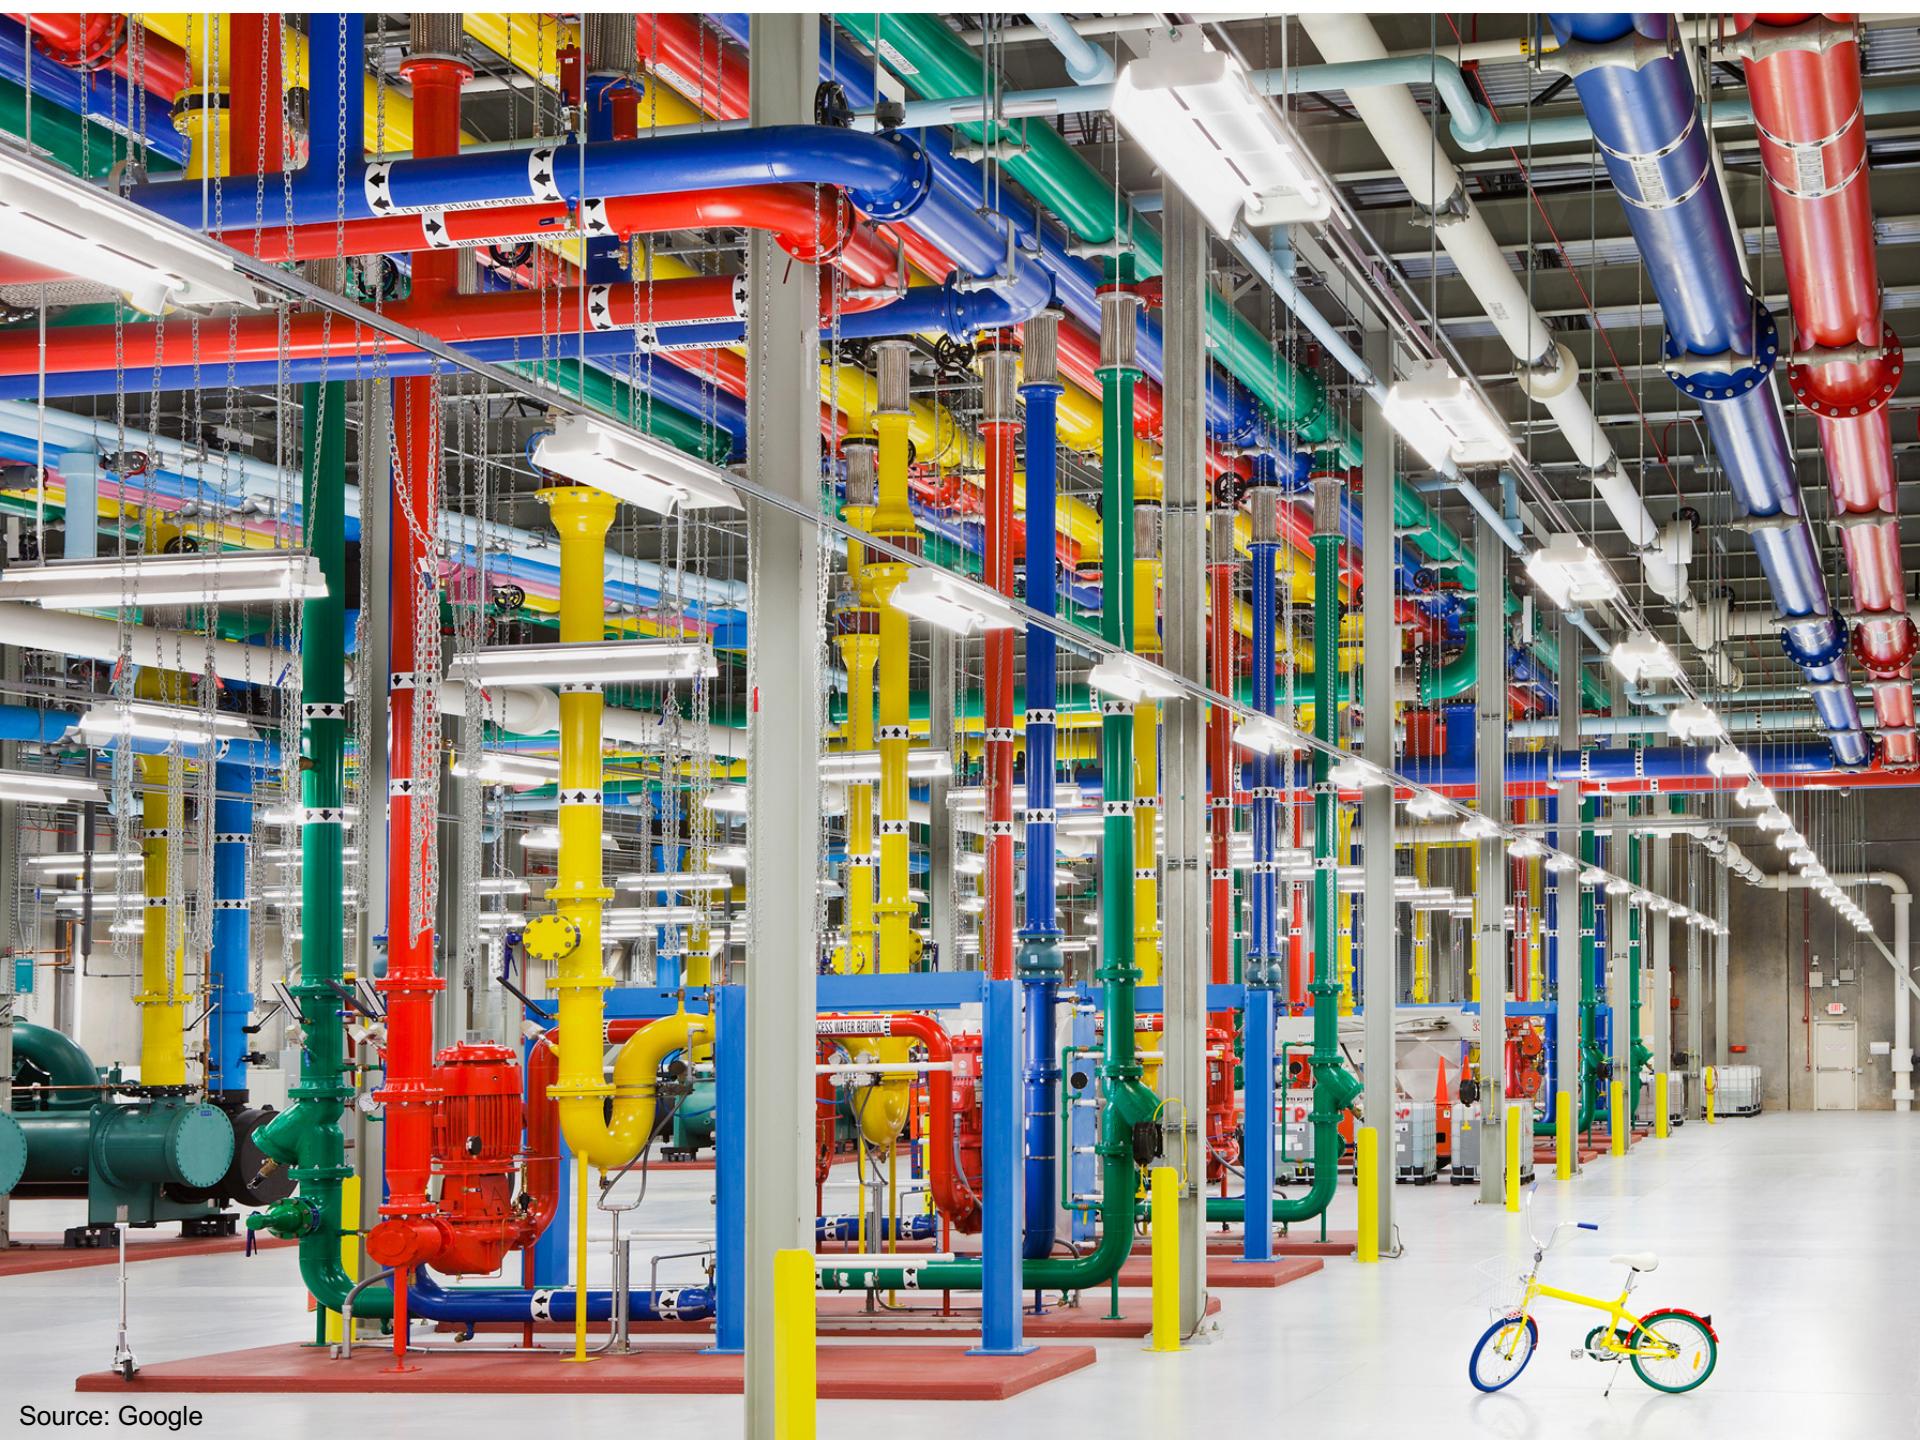


Datacenter cooling

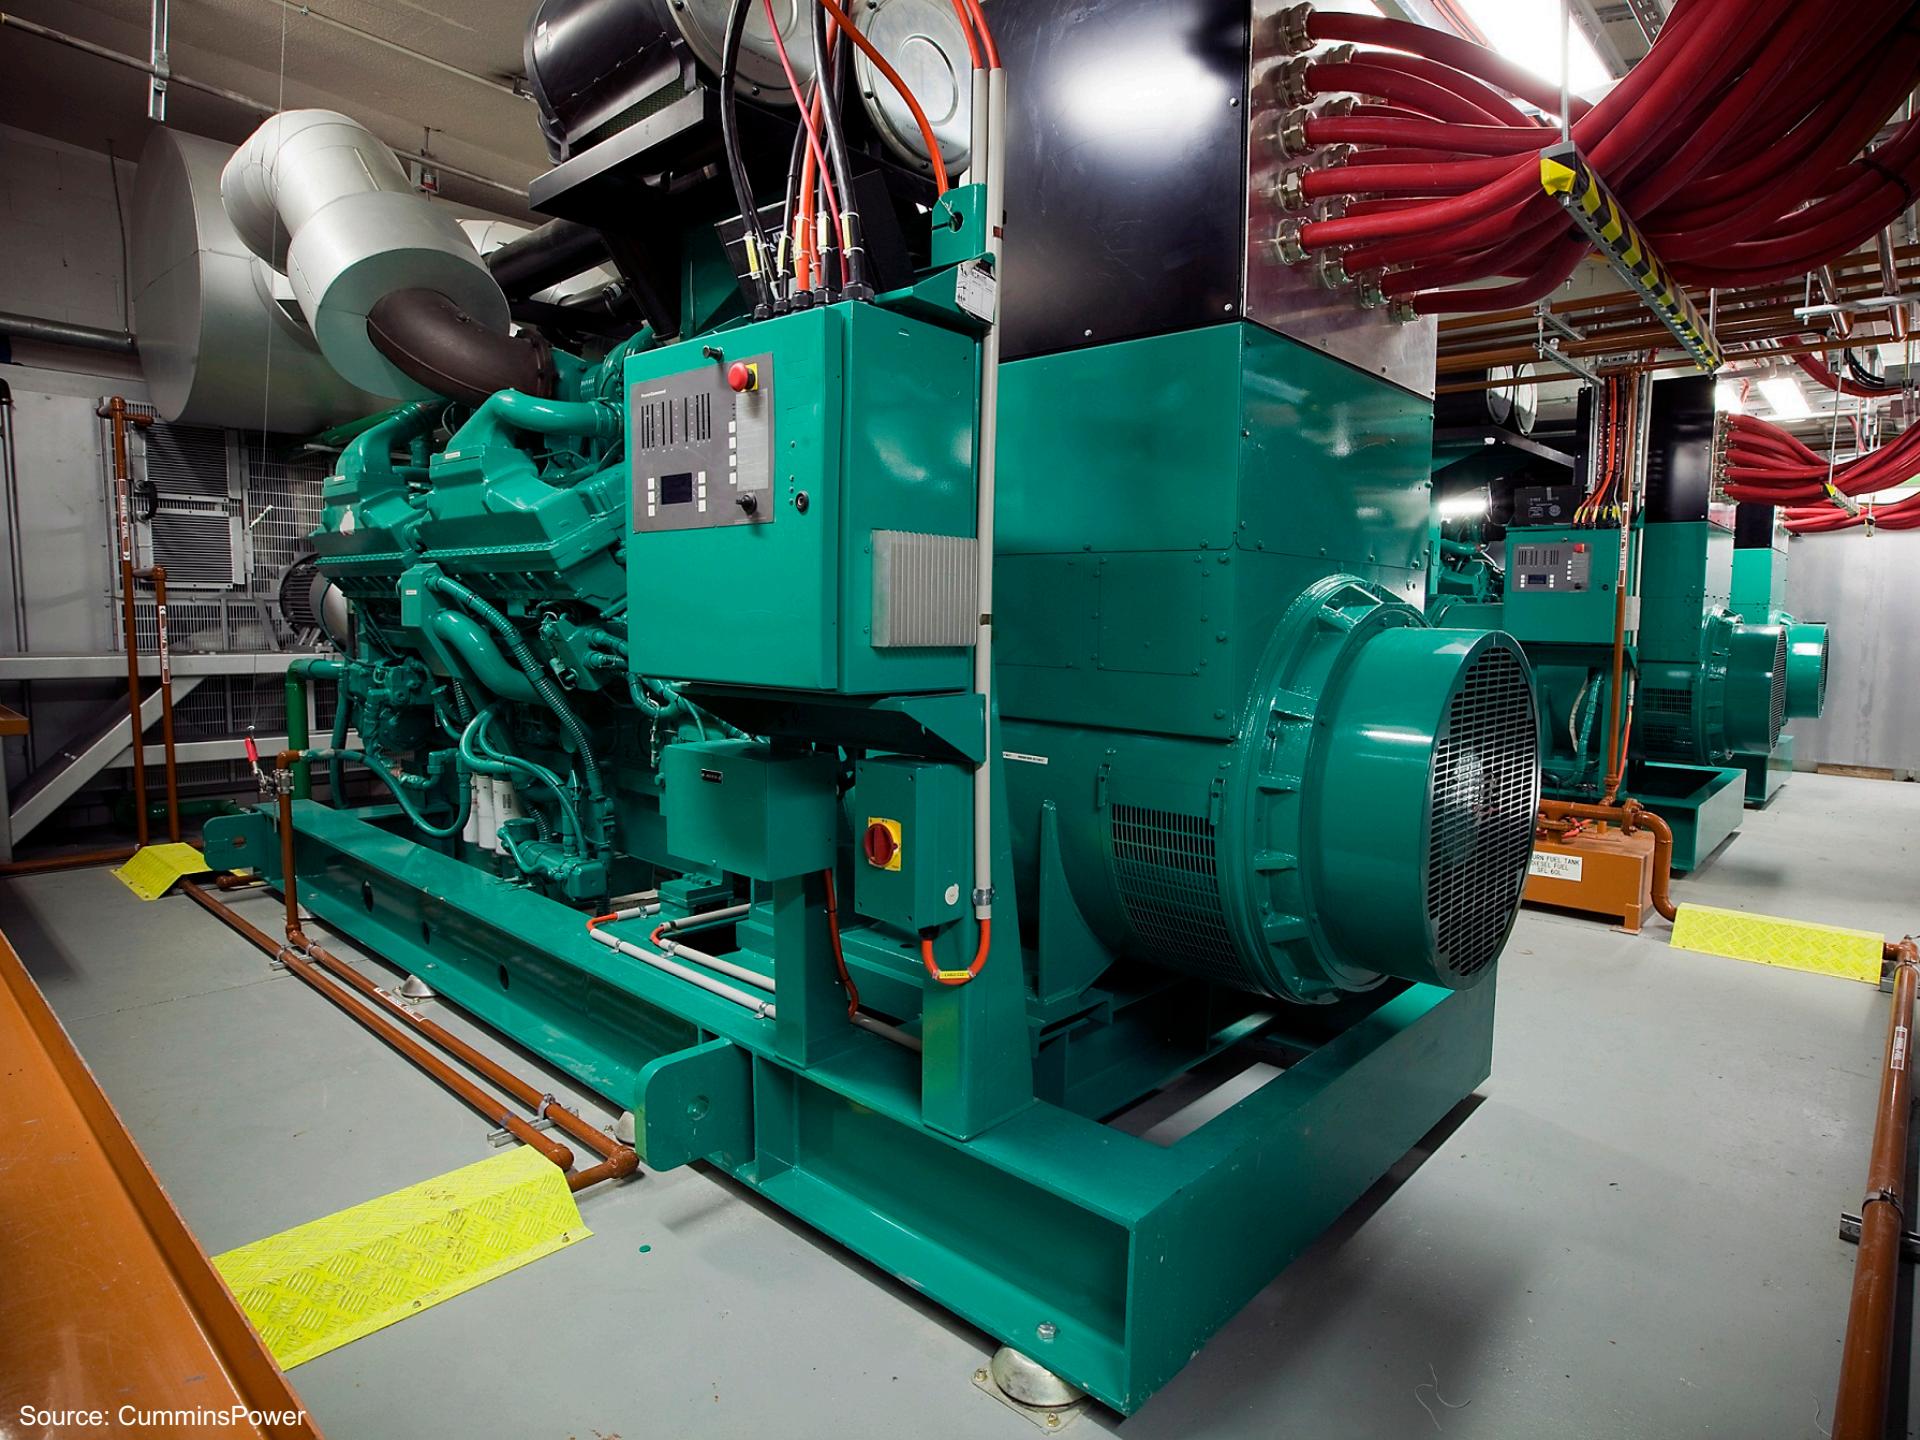
What's a computer?



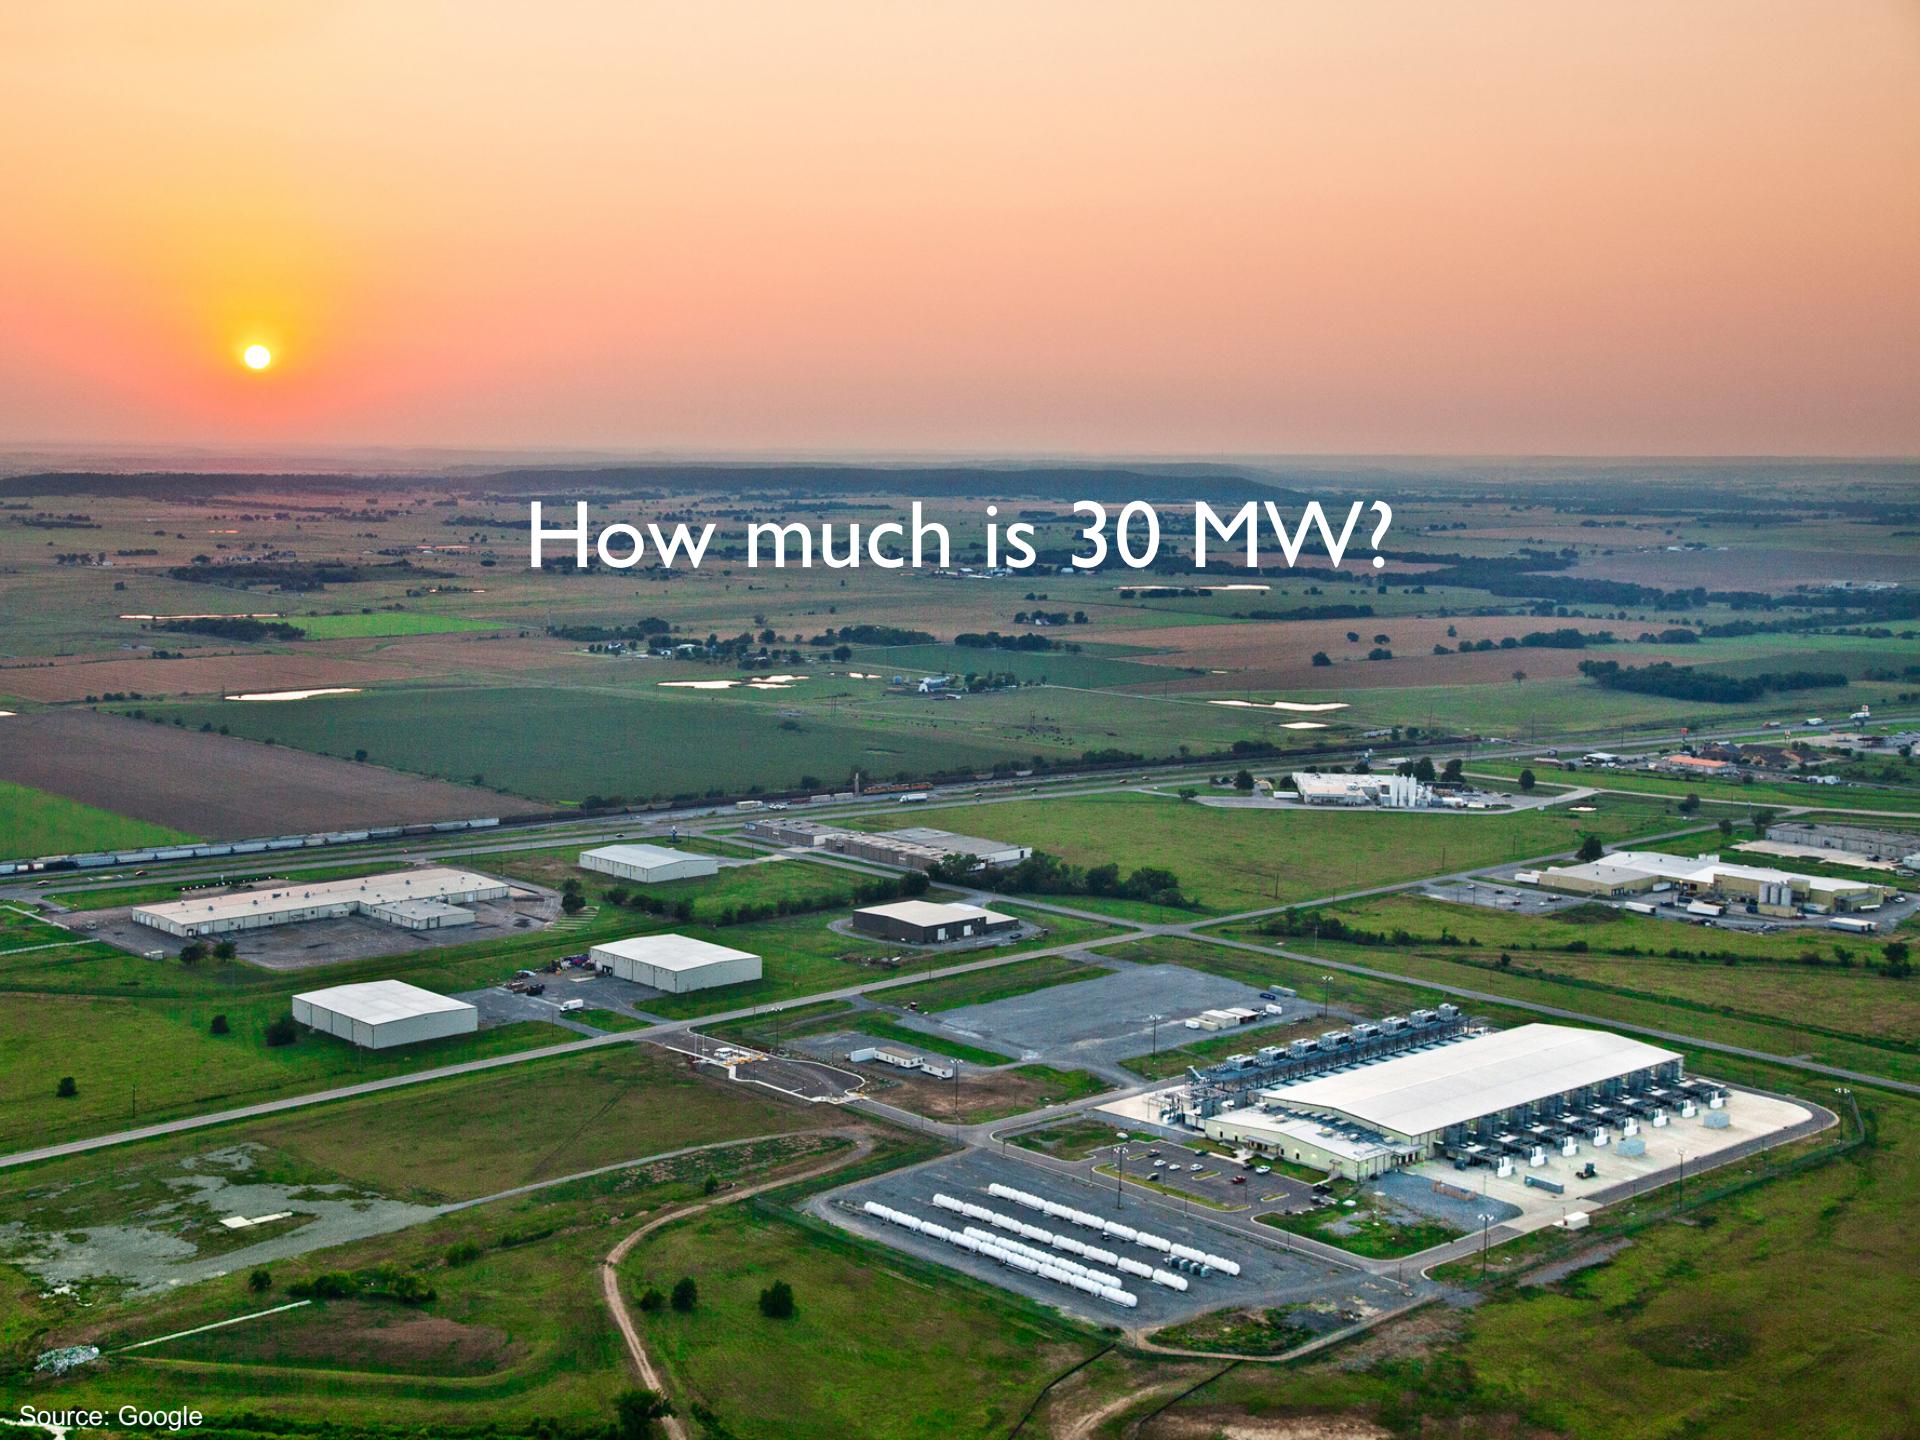




Source: Google

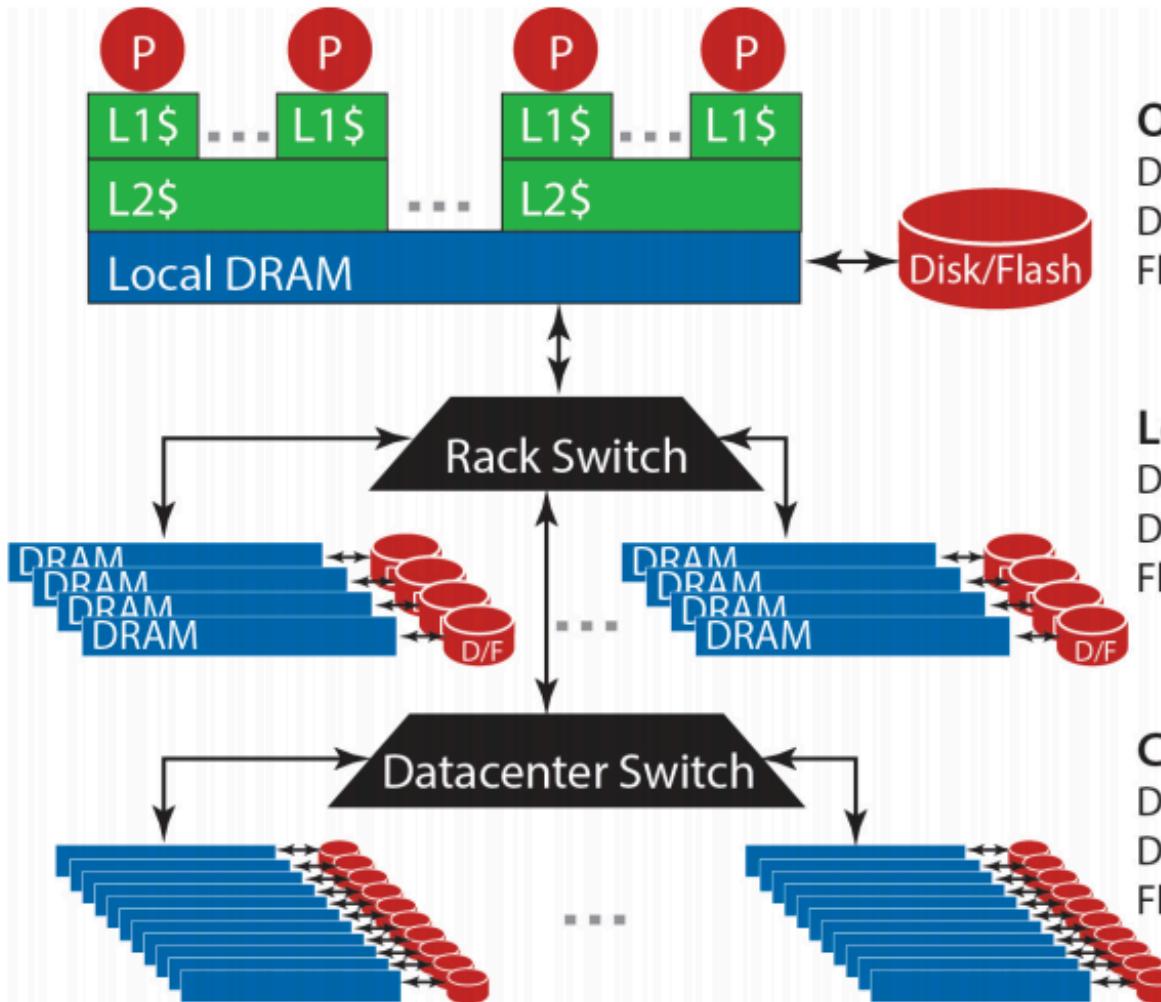






How much is 30 MW?

Datacenter Organization



One Server

DRAM: 16 GB, 100 ns, 20 GB/s
Disk: 2T B, 10 ms, 200 MB/s
Flash: 128 GB, 100 us, 1 GB/s

Local Rack (80 servers)

DRAM: 1 TB, 300 us, 100 MB/s
Disk: 160 TB, 11 ms, 100 MB/s
Flash: 20 TB, 400 us, 100 MB/s

Cluster (30 racks)

DRAM: 30 TB, 500 us, 10 MB/s
Disk: 4.80 PB, 12 ms, 10 MB/s
Flash: 600 TB, 600 us, 10 MB/s

The datacenter *is* the computer!

It's all about the right level of abstraction

Moving beyond the von Neumann architecture

What's the “instruction set” of the datacenter computer?

Hide system-level details from the developers

No more race conditions, lock contention, etc.

No need to explicitly worry about reliability, fault tolerance, etc.

Separating the *what* from the *how*

Developer specifies the computation that needs to be performed

Execution framework (“runtime”) handles actual execution

Wait, why do we care?

Mechanical Sympathy

“You don’t have to be an engineer to be
be a racing driver, but you do have to
have mechanical sympathy”
– Formula One driver Jackie Stewart



Data Science
Tools

Analytics
Infrastructure

Execution
Infrastructure

This Course

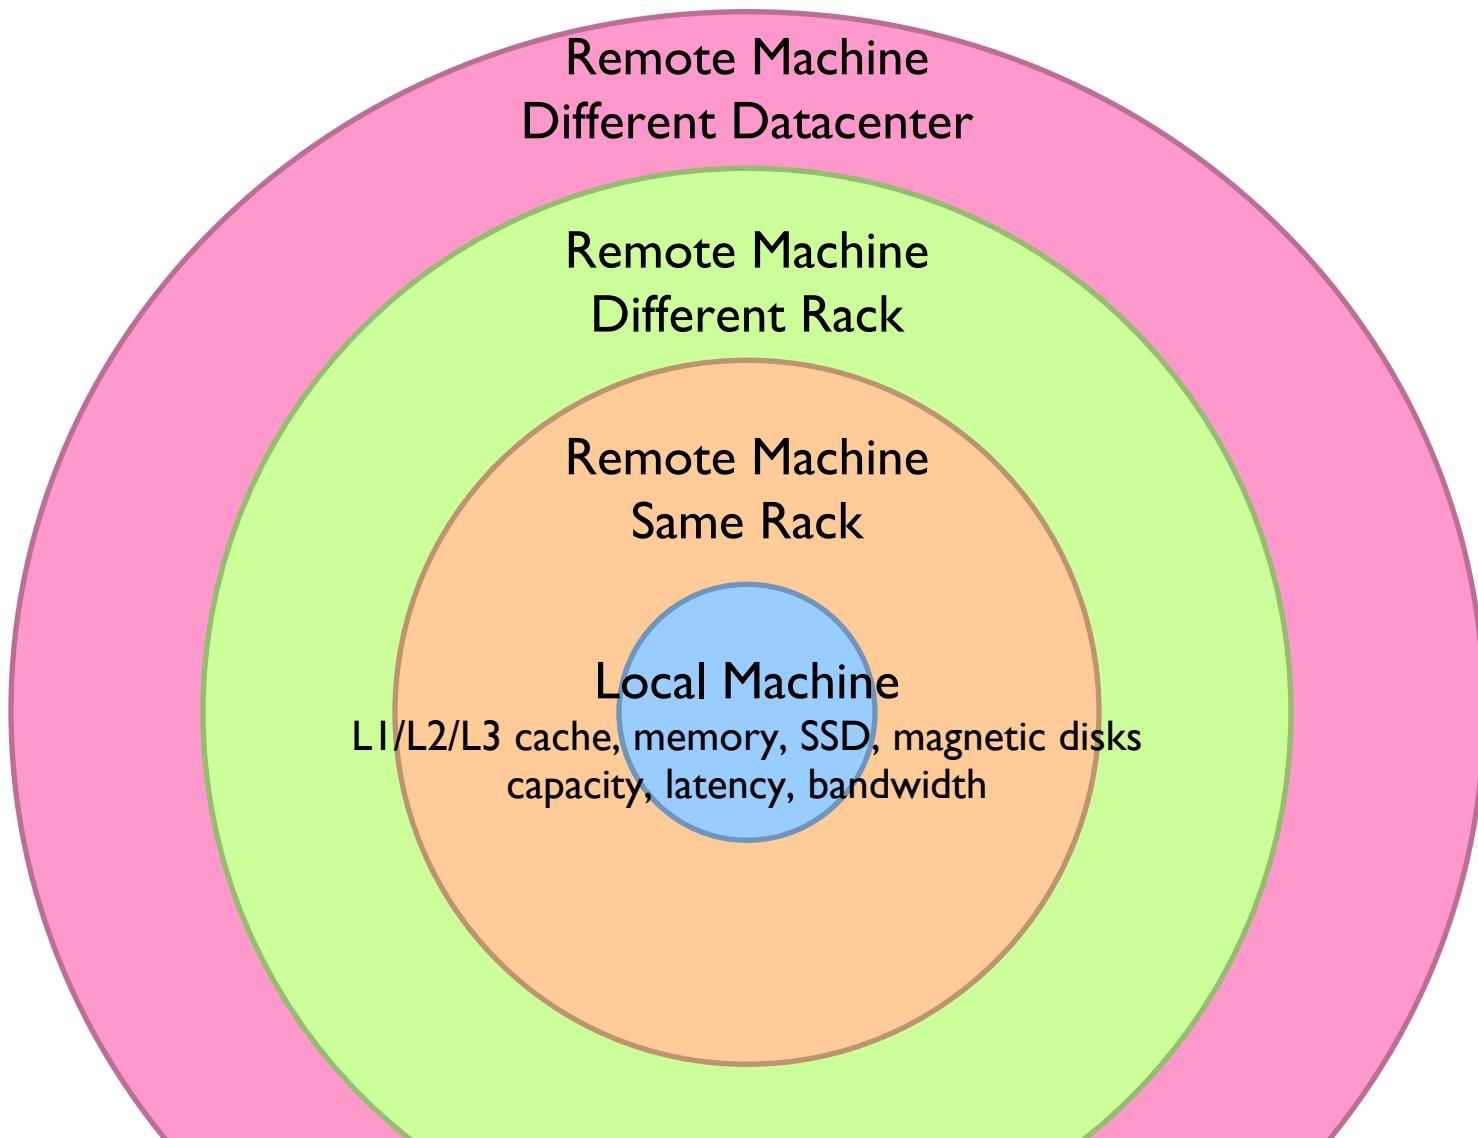
“big data stack”

Intuitions of time and space

How long does it take to read 100 TBs from 100 hard drives?
Now, what about SSDs?

How long will it take to exchange 1b key-value pairs:
Between machines on the same rack?
Between datacenters across the Atlantic?

Storage Hierarchy



Numbers Everyone Should Know

According to Jeff Dean

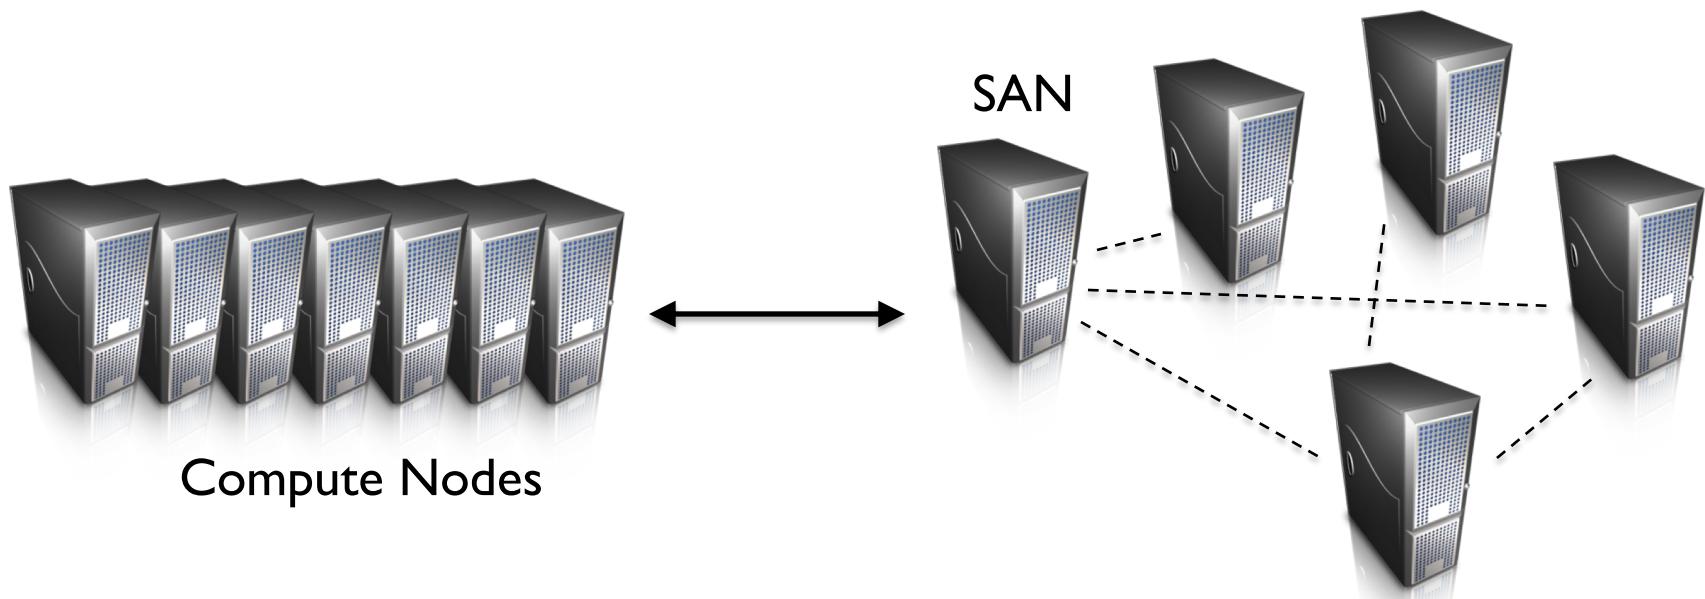
L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	10,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Hadoop Cluster Architecture



How do we get data to the workers?

Let's consider a typical supercomputer...



Sequoia will enable simulations that explore phenomena at a level of detail never before possible. Sequoia is dedicated to NNSA's Advanced Simulation and Computing (ASC) program for stewardship of the nation's nuclear weapons stockpile, a joint effort from LLNL, Los Alamos National Laboratory and Sandia National Laboratories.



Sequoia

16.32 PFLOPS

98,304 nodes with 1,572,864 million cores

1.6 petabytes of memory

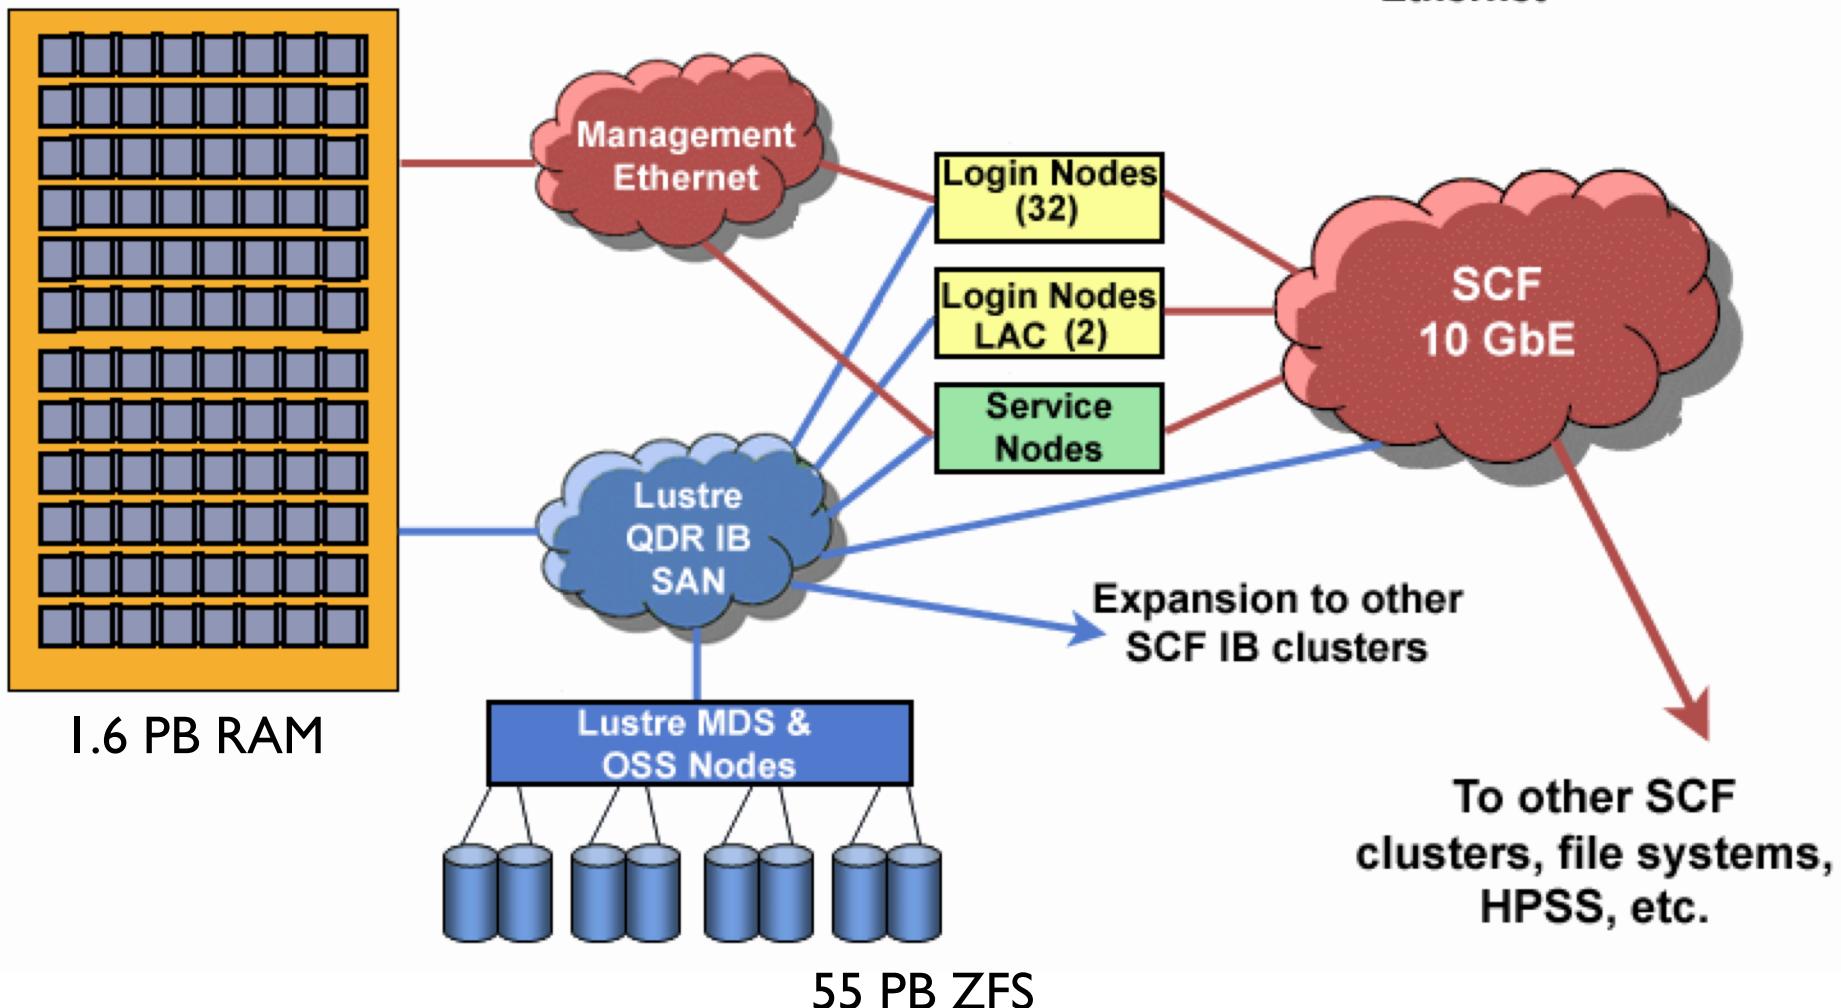
7.9 MWatts total power

Deployed in 2012, still #8 in TOP500 List (June 2018)

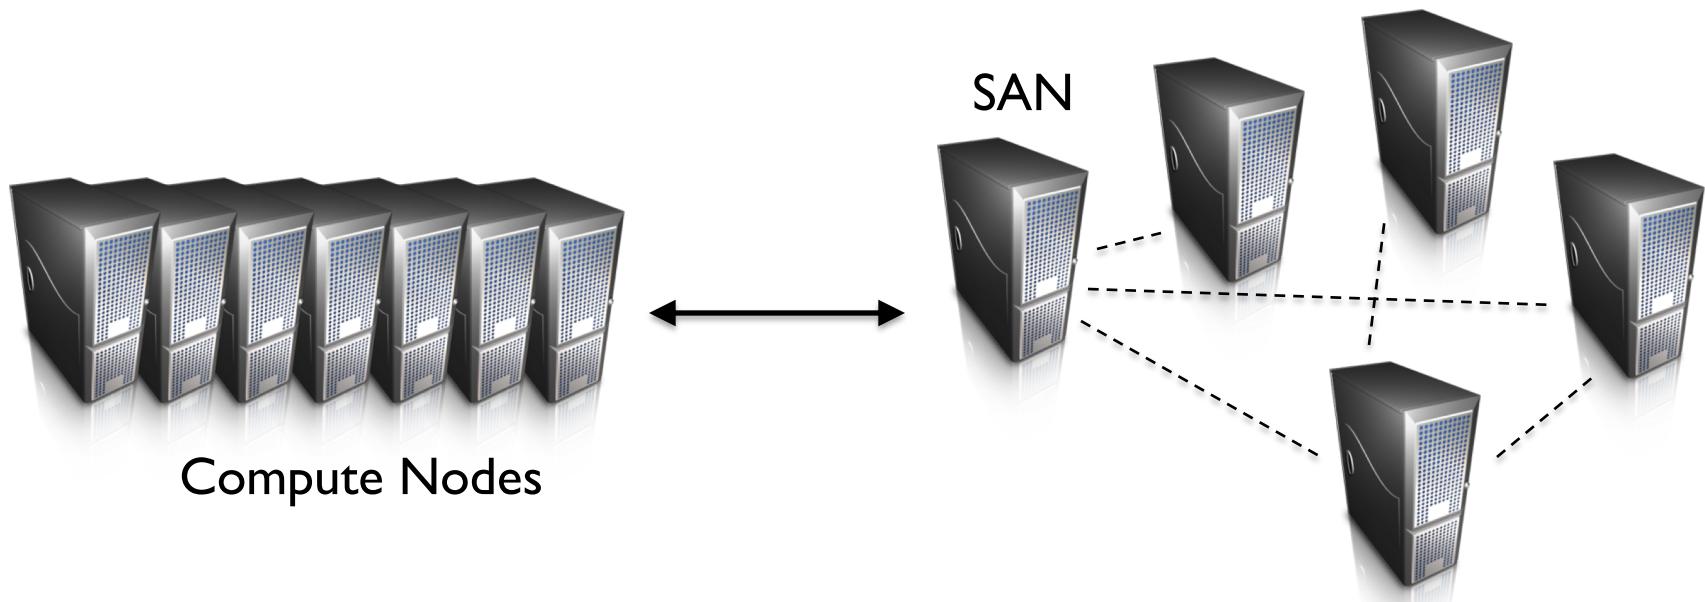
Sequoia

96 racks (12x8)
98,304 compute nodes
768 I/O nodes

- BG/Q 5D Torus Fabric
- QDR Infiniband
- Ethernet



Compute-Intensive vs. Data-Intensive

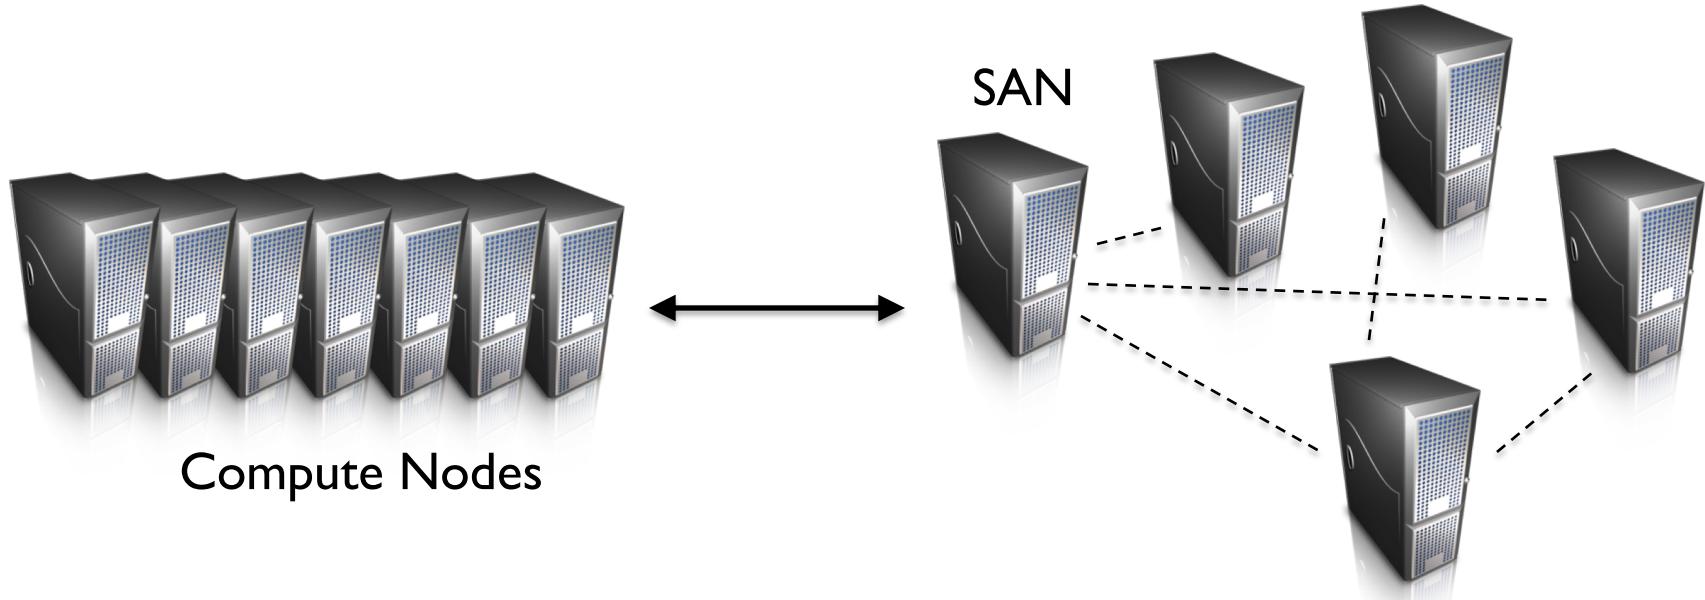


Why does this make sense for compute-intensive tasks?
What's the issue for data-intensive tasks?

What's the solution?

Don't move data to workers... move workers to the data!

Key idea: co-locate storage and compute
Start up worker on nodes that hold the data



What's the solution?

Don't move data to workers... move workers to the data!

Key idea: co-locate storage and compute

Start up worker on nodes that hold the data



We need a distributed file system for managing this

GFS (Google File System) for Google's MapReduce

HDFS (Hadoop Distributed File System) for Hadoop

GFS: Assumptions

Commodity hardware over “exotic” hardware

Scale “out”, not “up”

High component failure rates

Inexpensive commodity components fail all the time

“Modest” number of huge files

Multi-gigabyte files are common, if not encouraged

Files are write-once, mostly appended to

Logs are a common case

Large streaming reads over random access

Design for high sustained throughput over low latency

GFS: Design Decisions

Files stored as chunks
Fixed size (64MB)

Reliability through replication
Each chunk replicated across 3+ chunkservers

Single master to coordinate access and hold metadata
Simple centralized management

No data caching
Little benefit for streaming reads over large datasets

Simplify the API: not POSIX!
Push many issues onto the client (e.g., data layout)

HDFS = GFS clone (same basic ideas)

From GFS to HDFS

Terminology differences:

GFS master = Hadoop namenode

GFS chunkservers = Hadoop datanodes

Implementation differences:

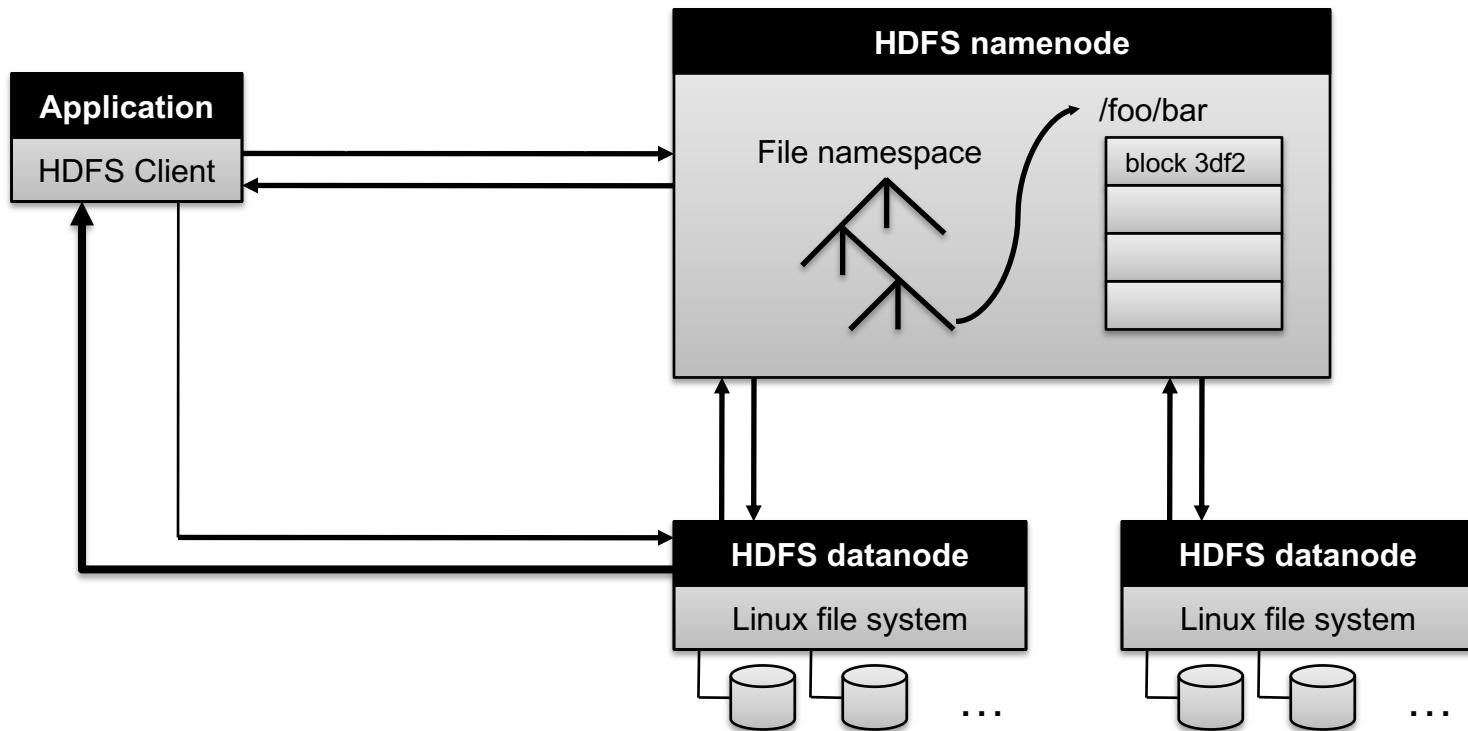
Different consistency model for file appends

Implementation language

Performance

For the most part, we'll use Hadoop terminology...

HDFS Architecture



Namenode Responsibilities

Managing the file system namespace

Holds file/directory structure, file-to-block mapping,
metadata (ownership, access permissions, etc.)

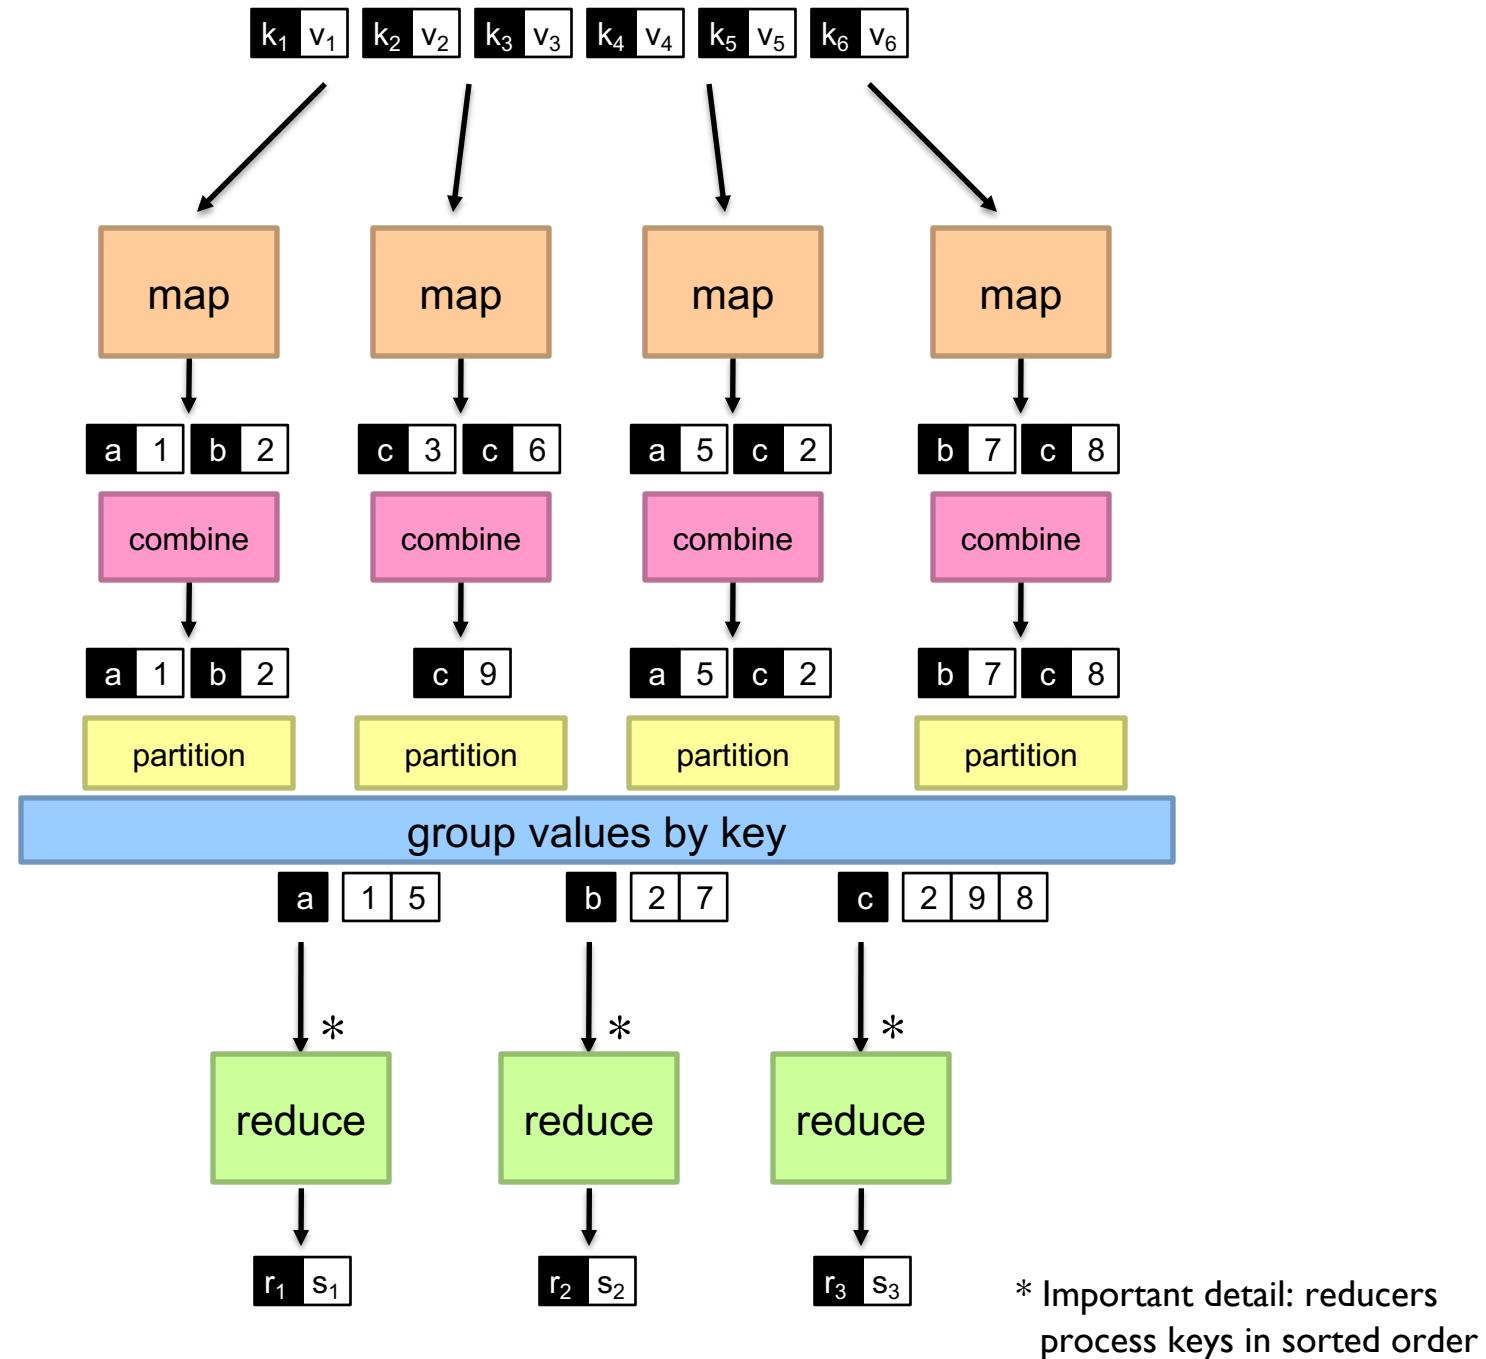
Coordinating file operations

Directs clients to datanodes for reads and writes
No data is moved through the namenode

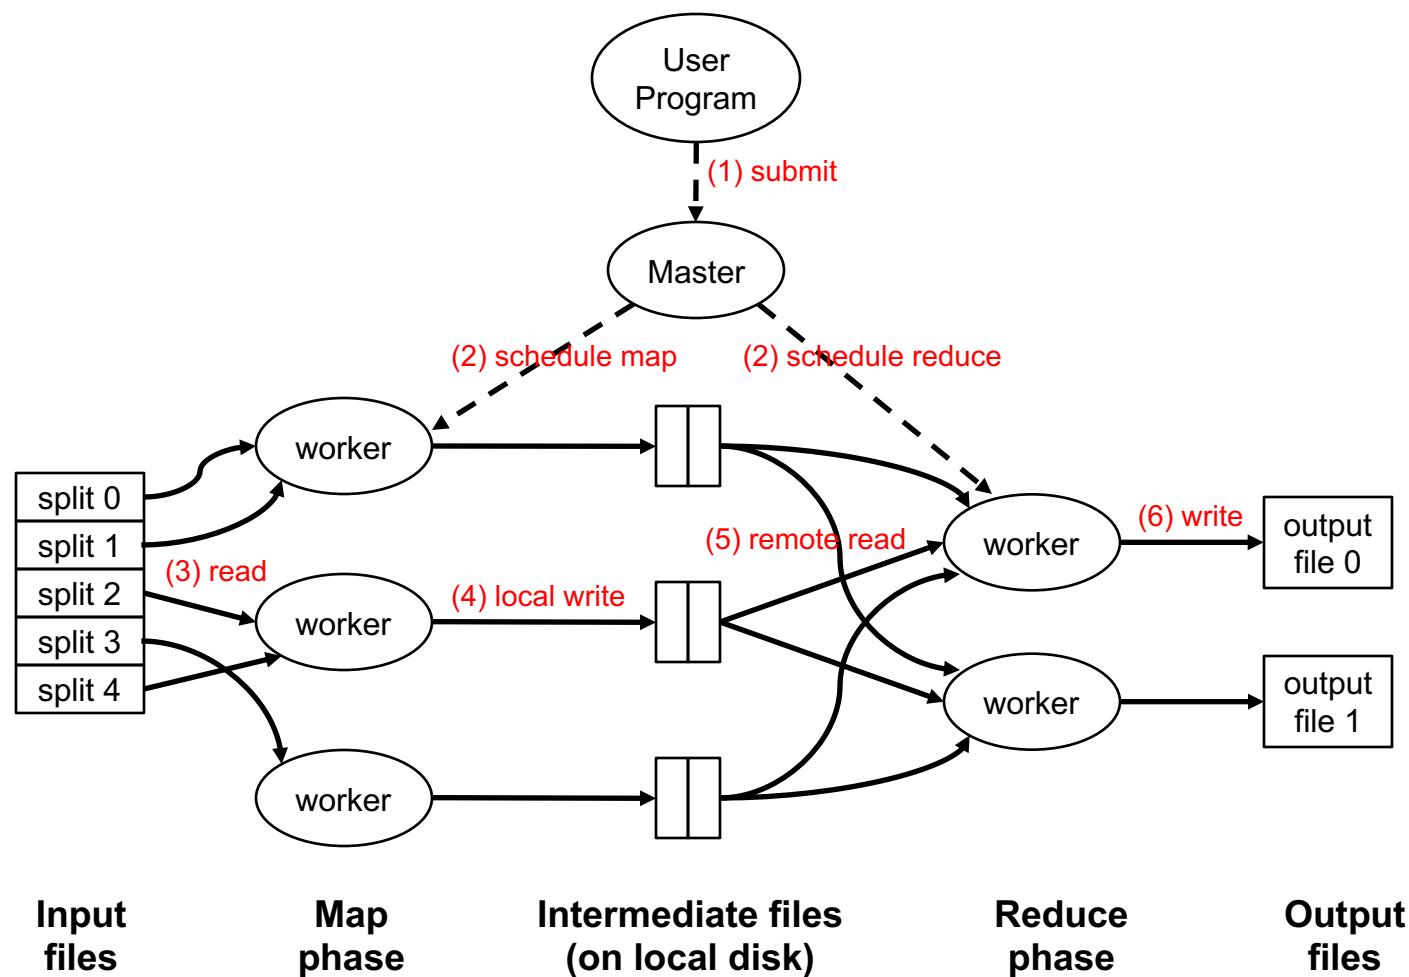
Maintaining overall health

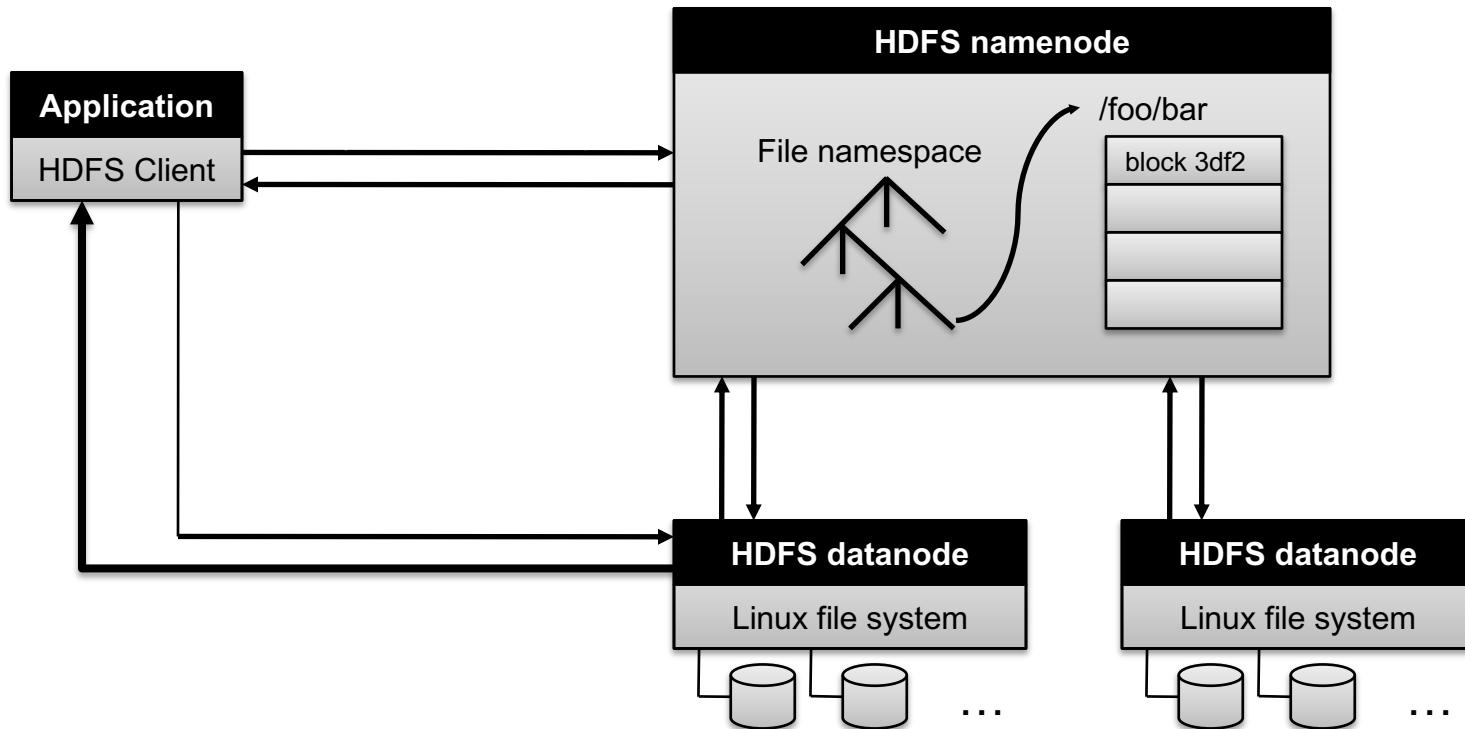
Periodic communication with the datanodes
Block re-replication and rebalancing
Garbage collection

Logical View

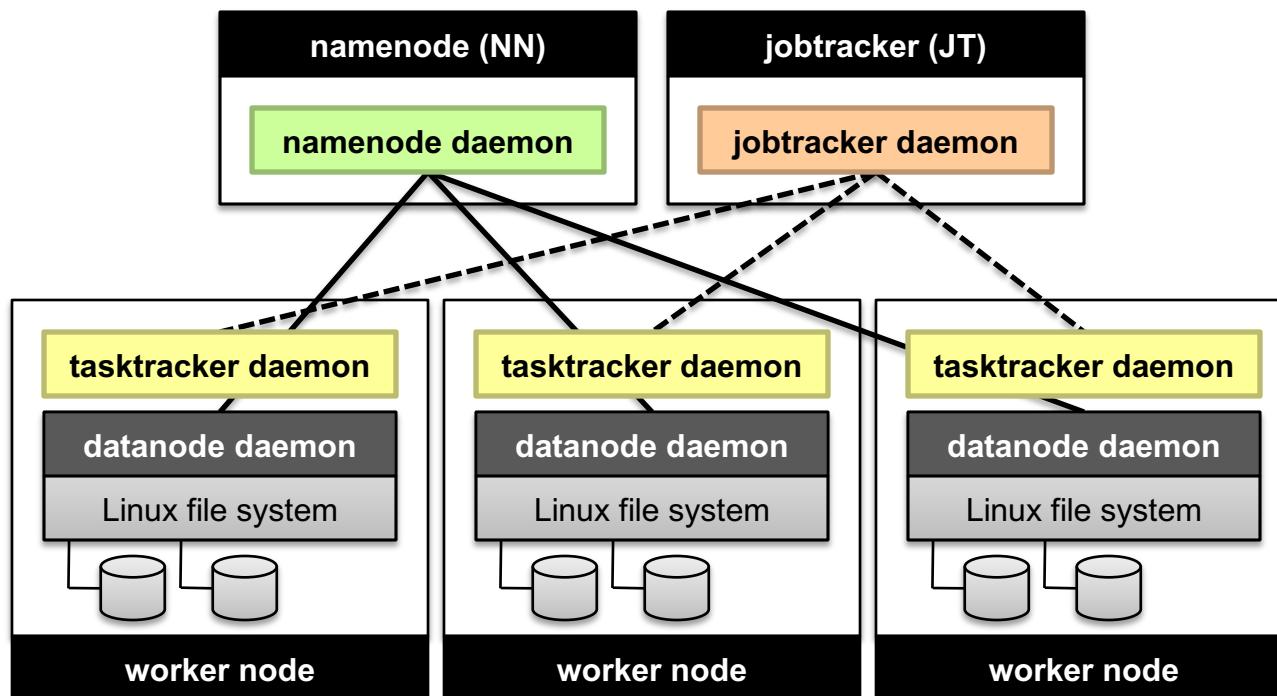


Physical View





Putting everything together...



Basic Cluster Components*

Namenode (NN)

Master for HDFS

Jobtracker (JT)

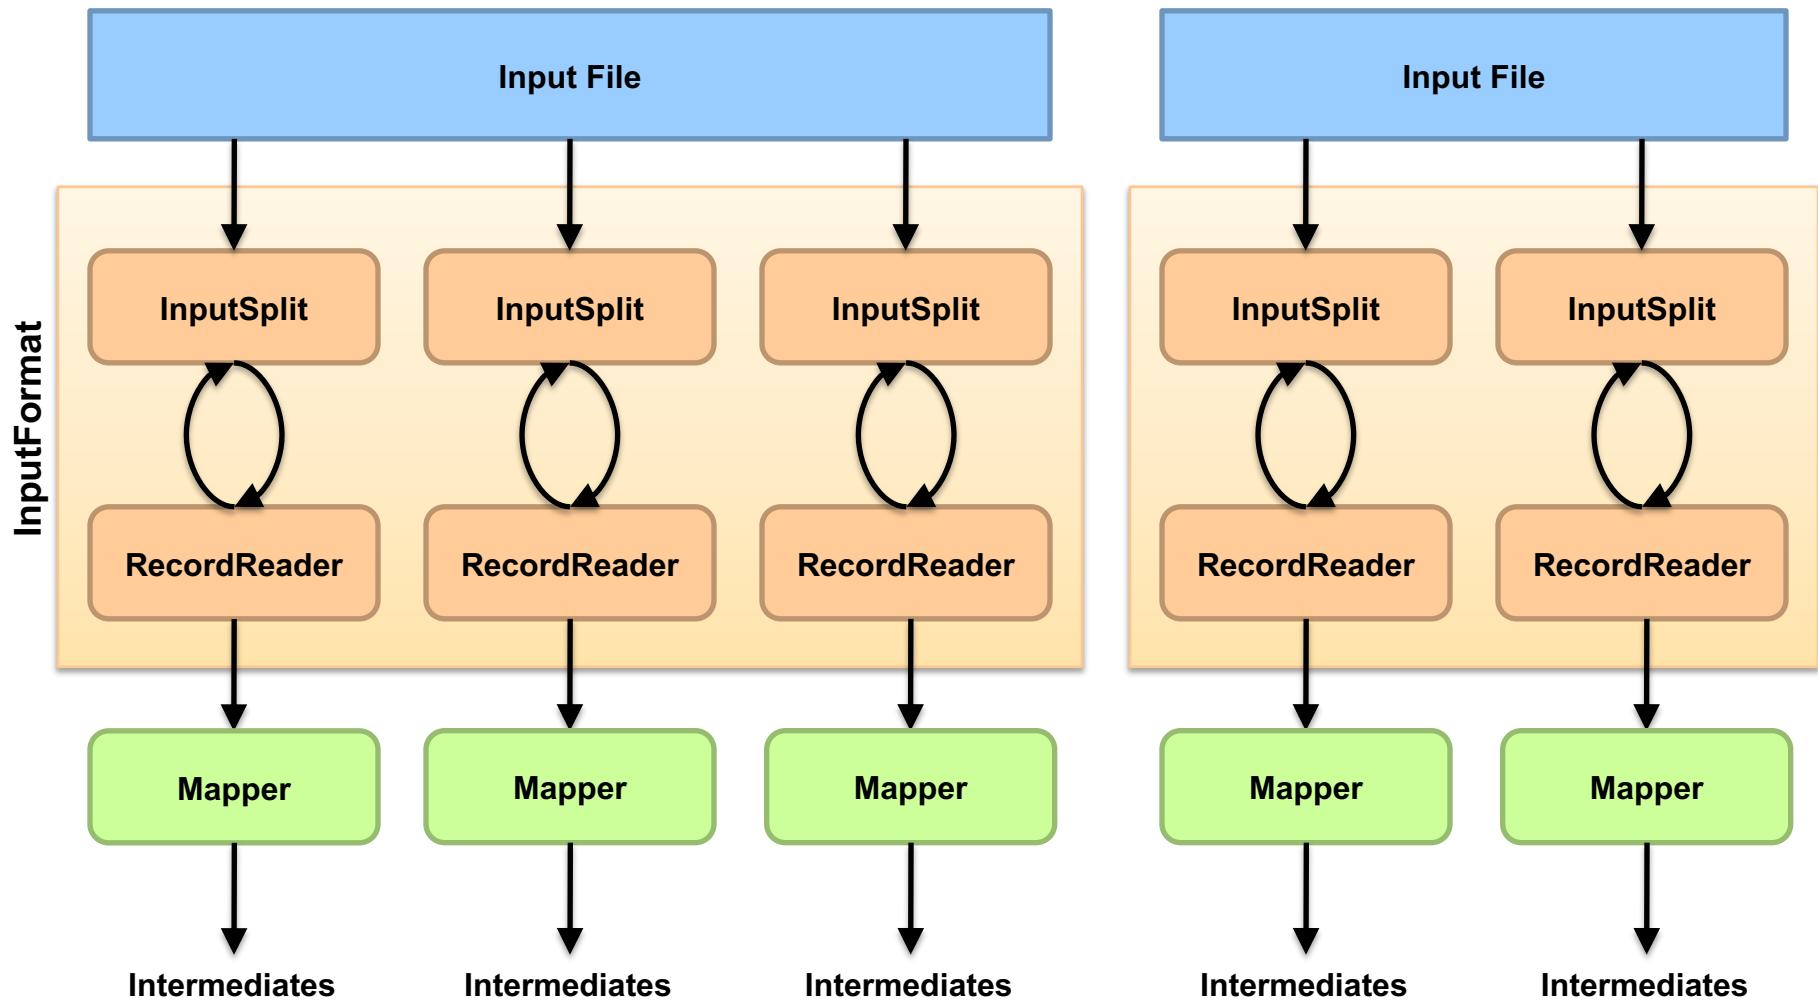
Coordinator for MapReduce jobs

On each of the worker machines:

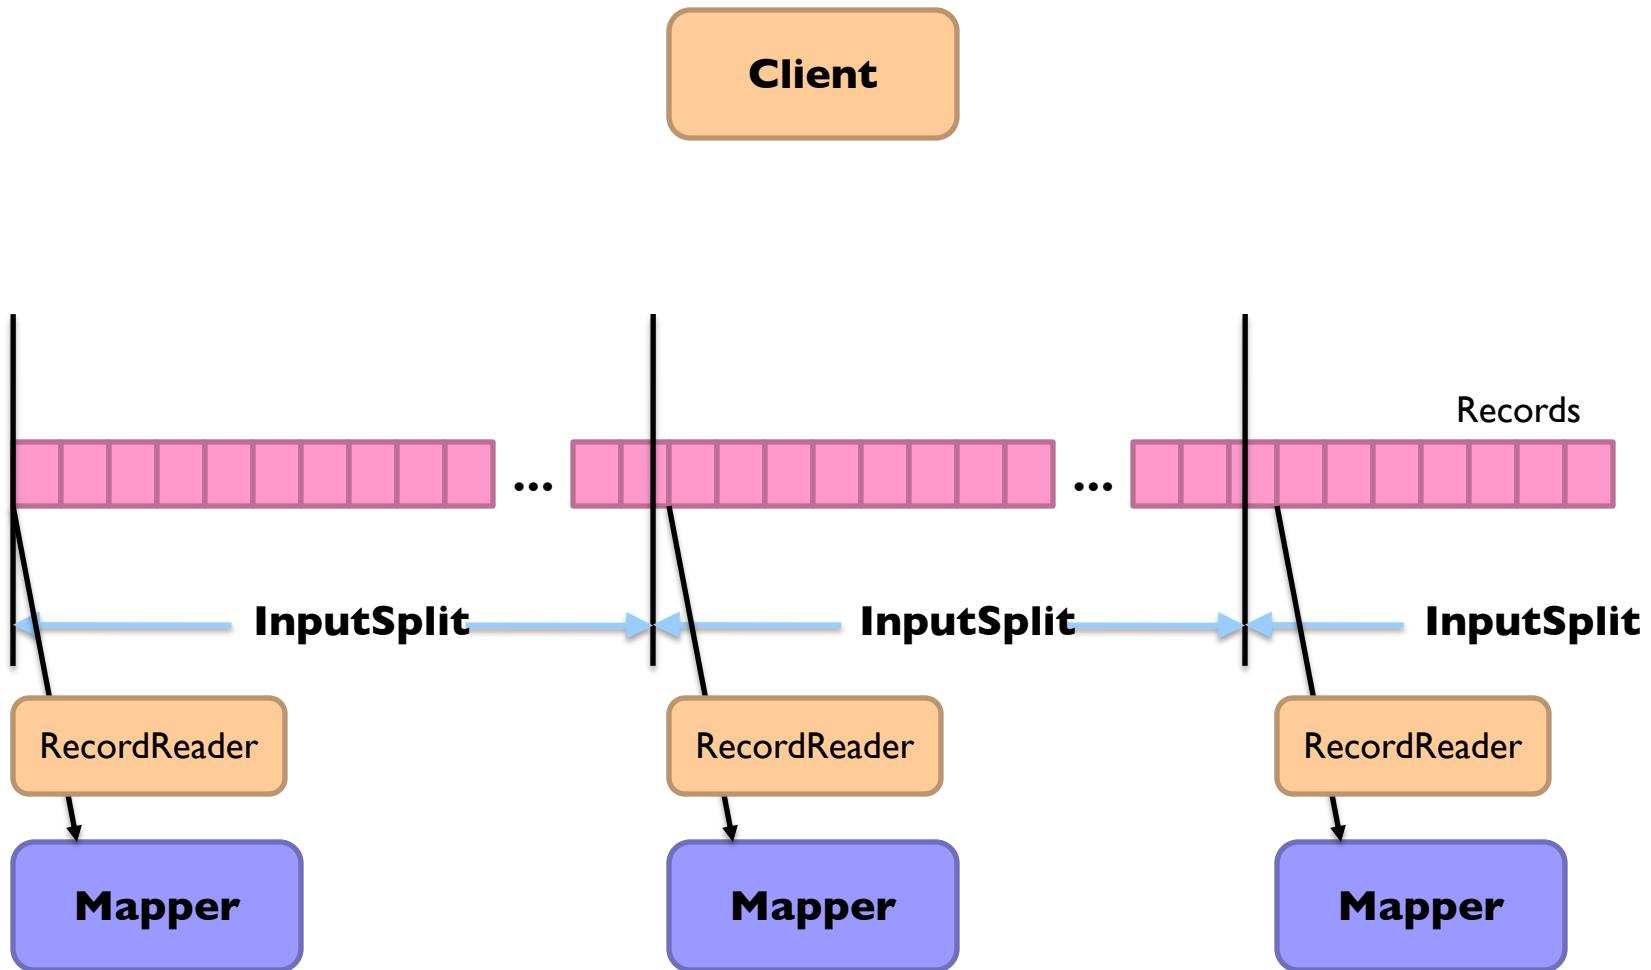
Tasktracker (TT): contains multiple task slots

Datanode (DN): serves HDFS data blocks

* Not quite... leaving aside YARN for now



What are these input split?



What are these input split?



Distributed Group By in MapReduce

Map side

Map outputs are buffered in memory in a circular buffer

When buffer reaches threshold, contents are “spilled” to disk

Spills are merged into a single, partitioned file (sorted within each partition)

Combiner runs during the merges

Reduce side

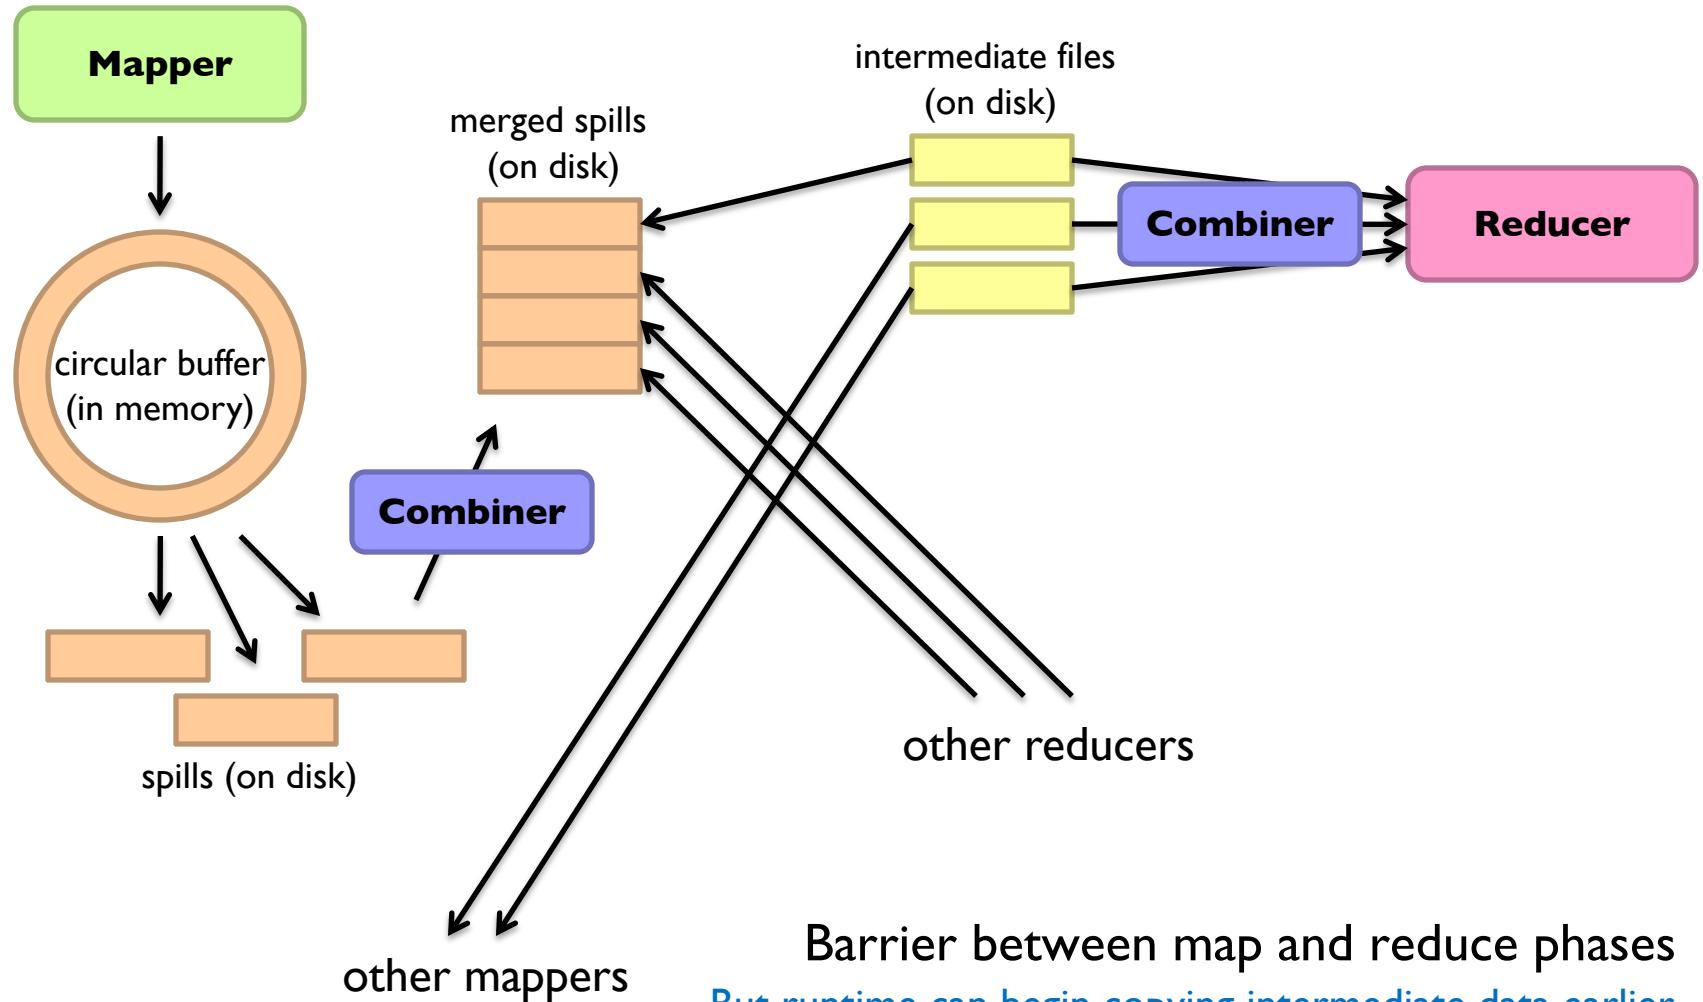
First, map outputs are copied over to reducer machine

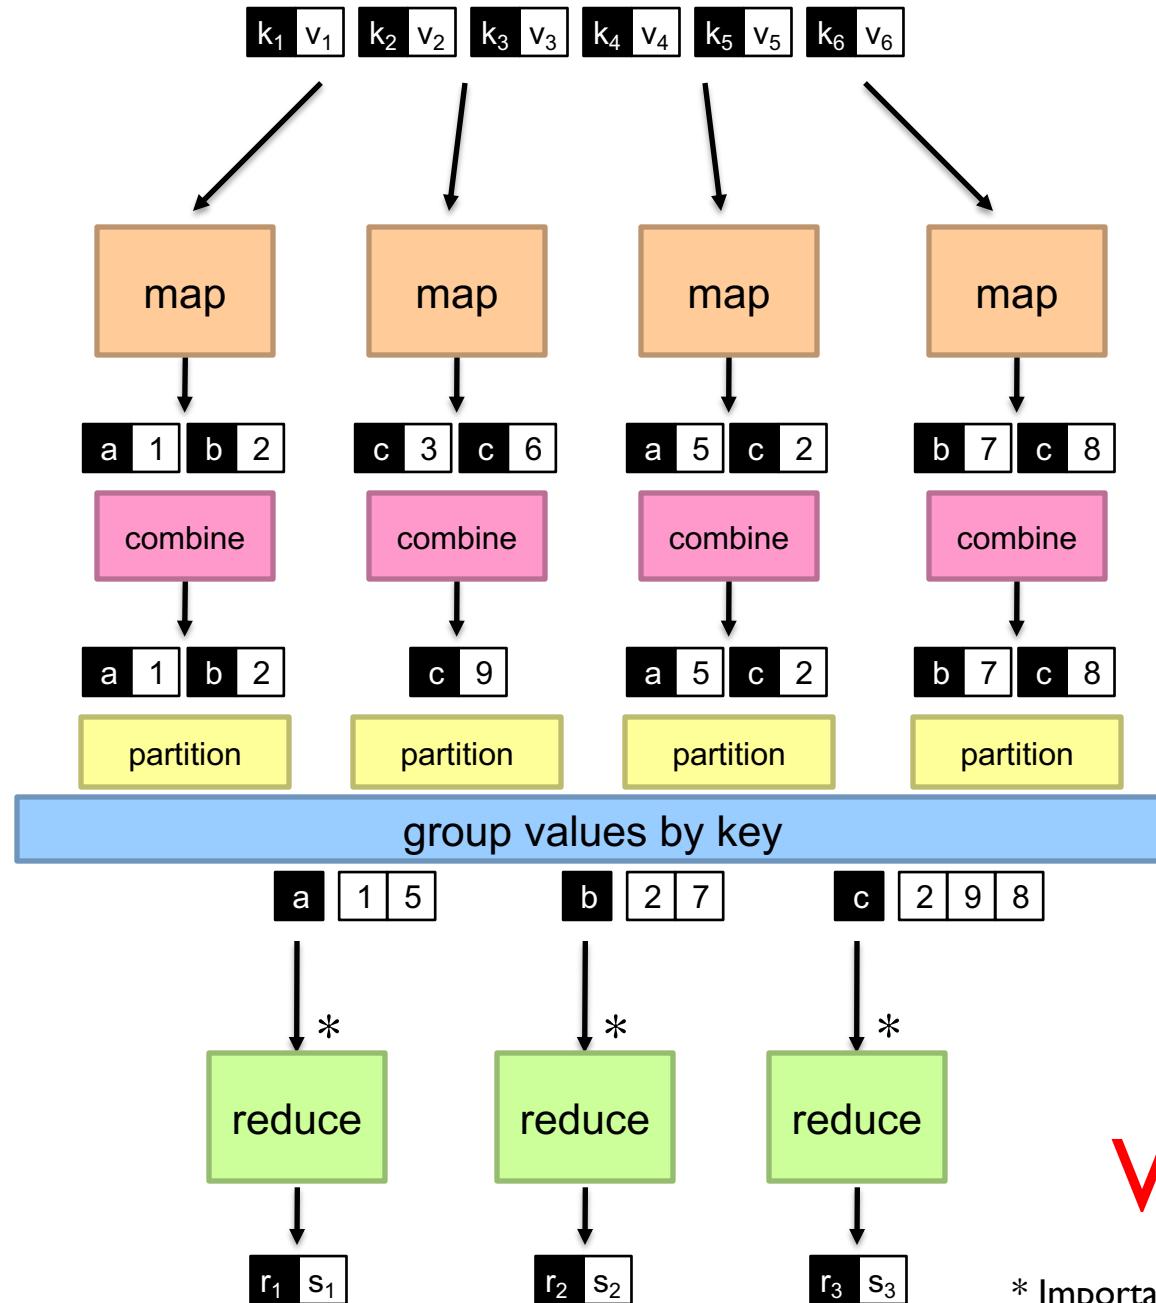
“Sort” is a multi-pass merge of map outputs (happens in memory and on disk)

Combiner runs during the merges

Final merge pass goes directly into reducer

Distributed Group By in MapReduce





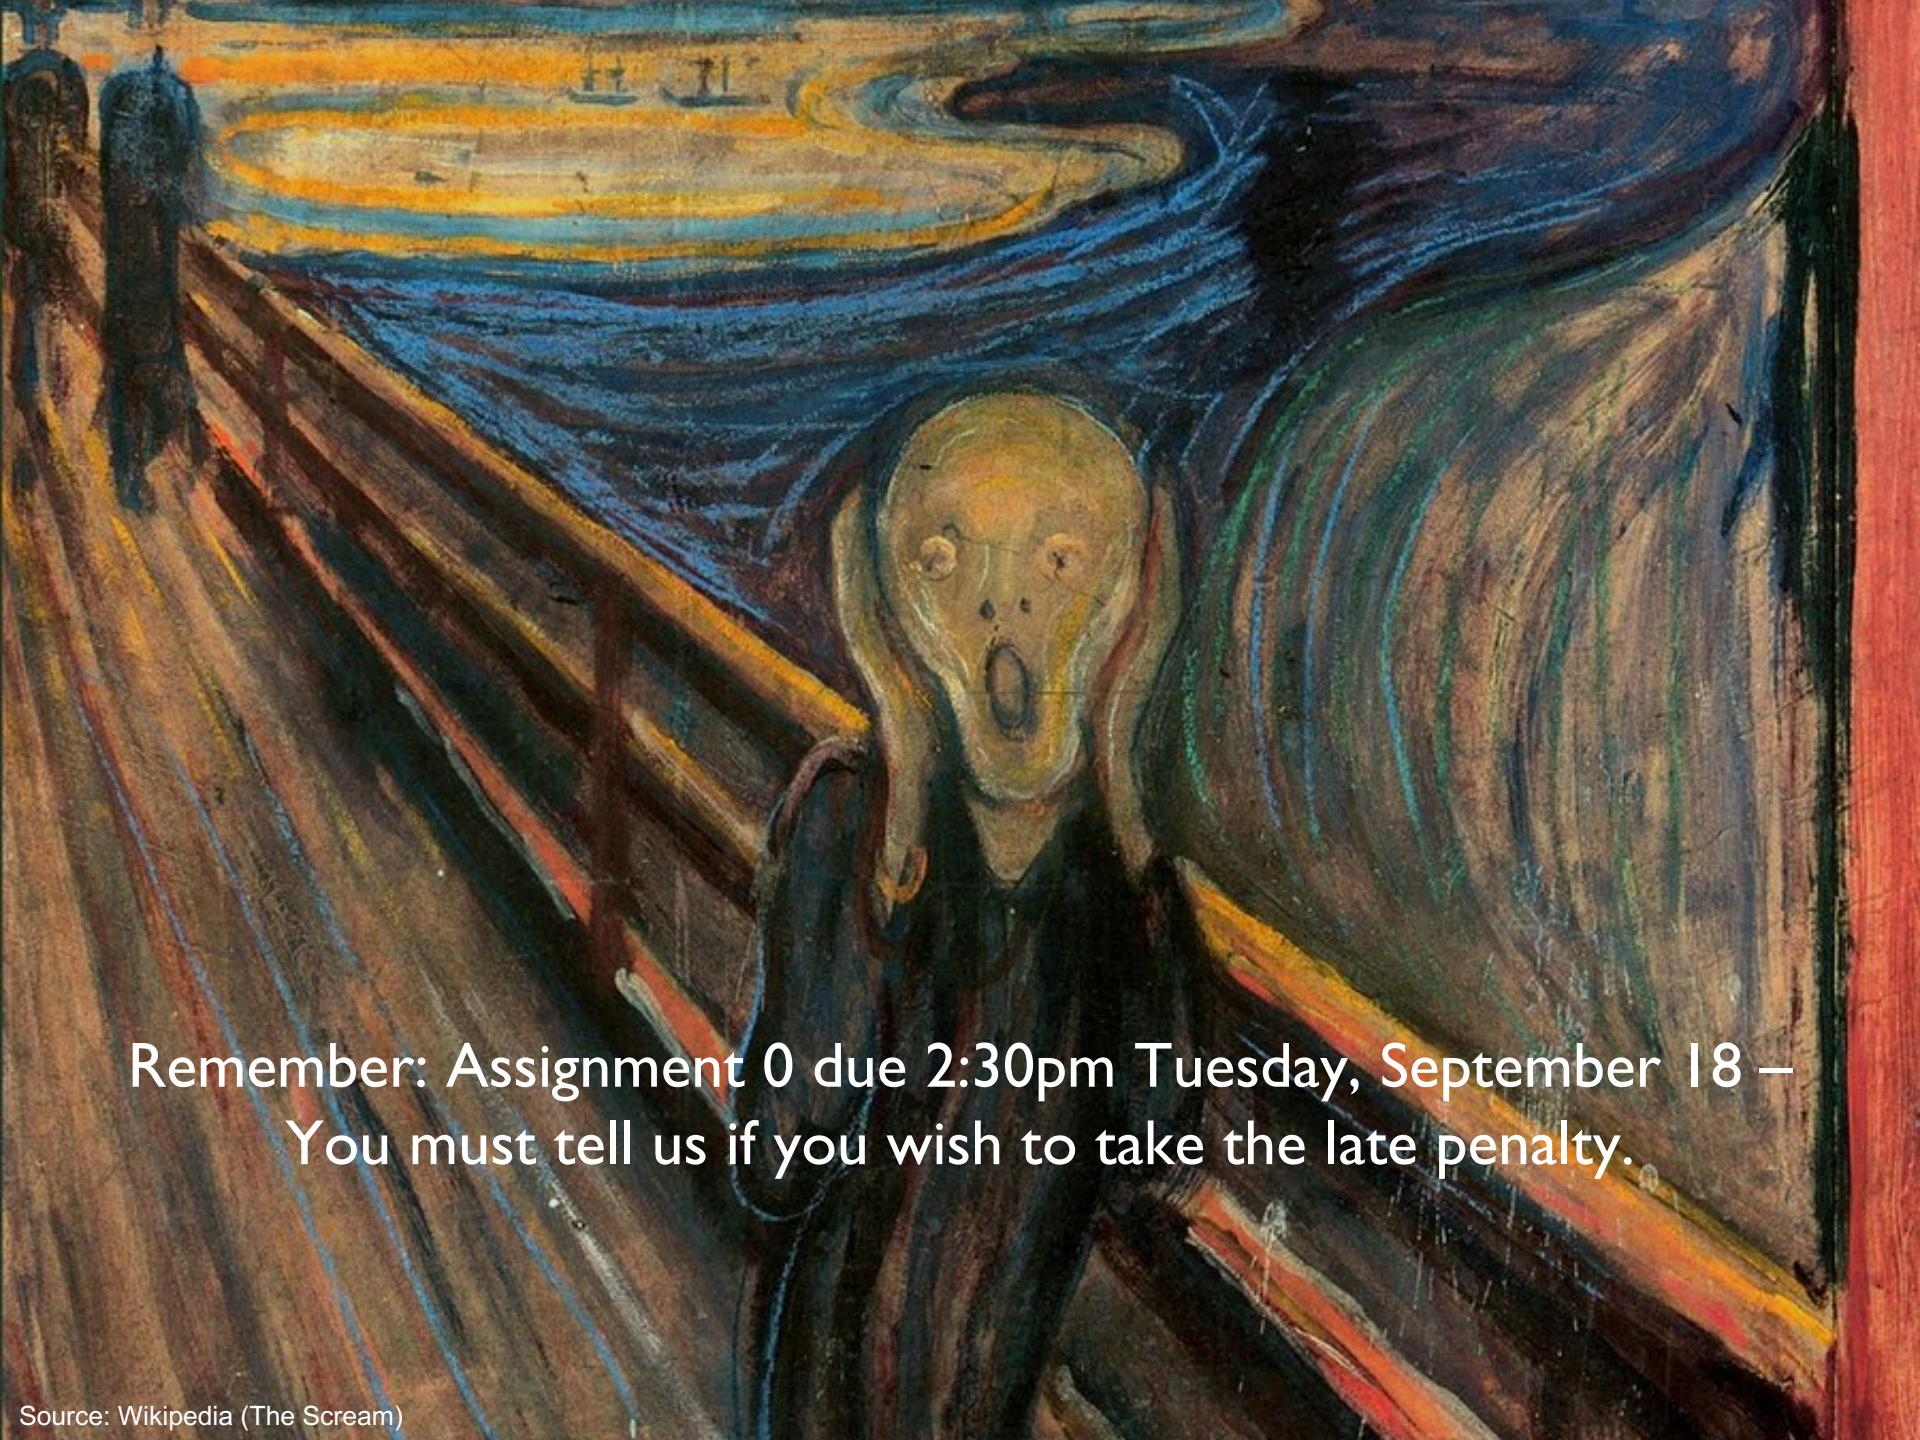
Why?

Law of Leaky Abstractions

All non-trivial abstractions, to some degree, are leaky.

Joel Spolsky

Remember logical vs. physical?



Remember: Assignment 0 due 2:30pm Tuesday, September 18 –
You must tell us if you wish to take the late penalty.